

## Scenario

In a distributed system with Byzantine fault tolerance, there are four nodes N1, N2, N3, and N4, where up to one node can be faulty ( $f = 1$ ). A client requests the values of keys  $x$  and  $y$  sequentially by sending two GET requests.

The system uses a quorum approach where each quorum requires at least three nodes ( $2f+1 = 3$ ) and ensures an overlap of at least two nodes between consecutive quorums.

Here's what happens during each request:

- For GET  $x$ , the client queries N1, N2, and N3.
  - Responses: N1 and N2 return  $x = 10$ , while N3 (a faulty node) returns  $x = 7$ .
- For GET  $y$ , the client queries N2, N3, and N4.
  - Responses: N2 and N4 return  $y = 20$ , while N3 (still faulty) returns  $y = 15$ .

## Question

Why is it important that there is an overlap of at least two nodes between the quorums for **GET  $x$**  and **GET  $y$** ? Describe how this overlap contributes to the reliability of the client's results despite node N3 being faulty.

## Answer

Since each quorum is of  $2f+1$  nodes and we can have  $f$  faulty nodes at most, PBFT nicely ensures that we always have  $f+1$  non-faulty nodes. For any configuration of nodes that is valid according to PBFT rules, this also means that we will have an overlap of  $f+1$  nodes across requests. The overlap ensures that at least one non-faulty node is common across both requests, which means there will be consistent information despite N3's faulty behaviour. It also allows the client to cross-validate data across requests, ensuring that any discrepancies between responses can be attributed to the faulty node, not to inconsistencies in the system.

That said, the  $2f+1$  quorum already ensures the reliability of results returned by a majority, so this is an additional reassurance at best.