

CS 582: Distributed Systems

Remote Procedure Calls

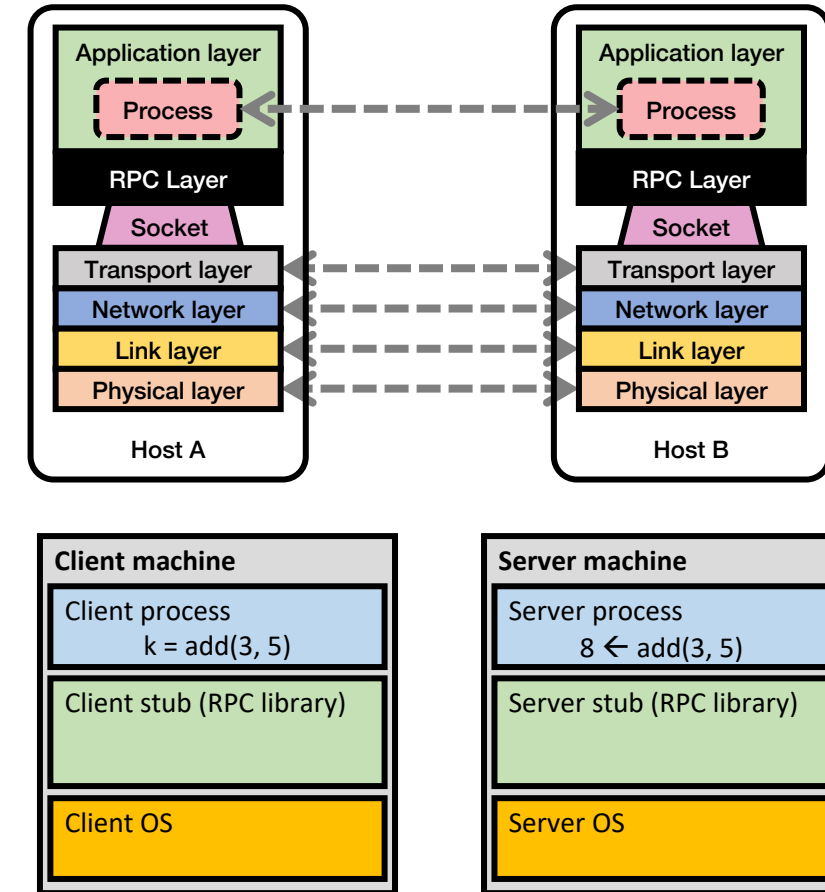


Dr. Zafar Ayyub Qazi

Fall 2024

Remote Procedure Call (RPC)

- Why RPC?
- What are RPCs?
- Issues in implementing RPCs
- How do RPCs work?



Learning Outcomes: RPCs


You should be able to:

- ❑ Define Remote Procedure Calls (RPCs) and explain their purpose in distributed systems
- ❑ Describe the challenges associated with implementing remote procedure calls
- ❑ Compare and contrast local and remote procedure calls, evaluating their respective advantages and disadvantages
- ❑ Analyze how the choice of transport protocol impacts RPC performance and delays
- ❑ Examine how failures affect RPC semantics and predict potential outcomes in failure scenarios
- ❑ Evaluate different strategies for handling failures in RPCs and propose appropriate solutions for given scenarios


Everyone uses RPCs

- Google gRPC
- Facebook/Apache Thrift
- Twitter Finagle
- CS 582 assignments
- ...

Why RPCs? Or whats wrong with sockets?



When designing a distributed system often there is a desire to make it “work like a single system”



Principle of Transparency

- Principle of transparency: Hide the fact that resources are physically distributed across multiple computers
 - Access resources the same way as we would do locally
 - Users can't tell where resources are physically located

Simple socket-based program


```
if ((sockfd = socket (AF_INET, SOCK_STREAM, 0)) < 0) {  
    perror("Socket creation");  
    exit(2);  
}
```

```
memset(&servaddr, 0, sizeof(servaddr));  
servaddr.sin_family = AF_INET;  
servaddr.sin_addr.s_addr = inet_addr(argv[1]);  
servaddr.sin_port = htons(SERV_PORT); // to big-endian
```

```
if (connect(sockfd, (struct sockaddr *) &servaddr,  
            sizeof(servaddr)) < 0) {  
    perror("Connect to server");  
    exit(3);  
}
```

```
send(sockfd, buf, strlen(buf), 0);
```

```
// Create a socket for the client
if ((sockfd = socket (AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("Socket creation");
    exit(2);
}
```

```
// Set server address and port
memset(&servaddr, 0, sizeof(servaddr));
servaddr.sin_family = AF_INET;
servaddr.sin_addr.s_addr = inet_addr(argv[1]);
servaddr.sin_port = htons(SERV_PORT); // to big-endian
```

```
// Establish TCP connection
if (connect(sockfd, (struct sockaddr *) &servaddr,
            sizeof(servaddr)) < 0) {
    perror("Connect to server");
    exit(3);
}
```

```
// Transmit the data over the TCP connection
send(sockfd, buf, strlen(buf), 0);
```

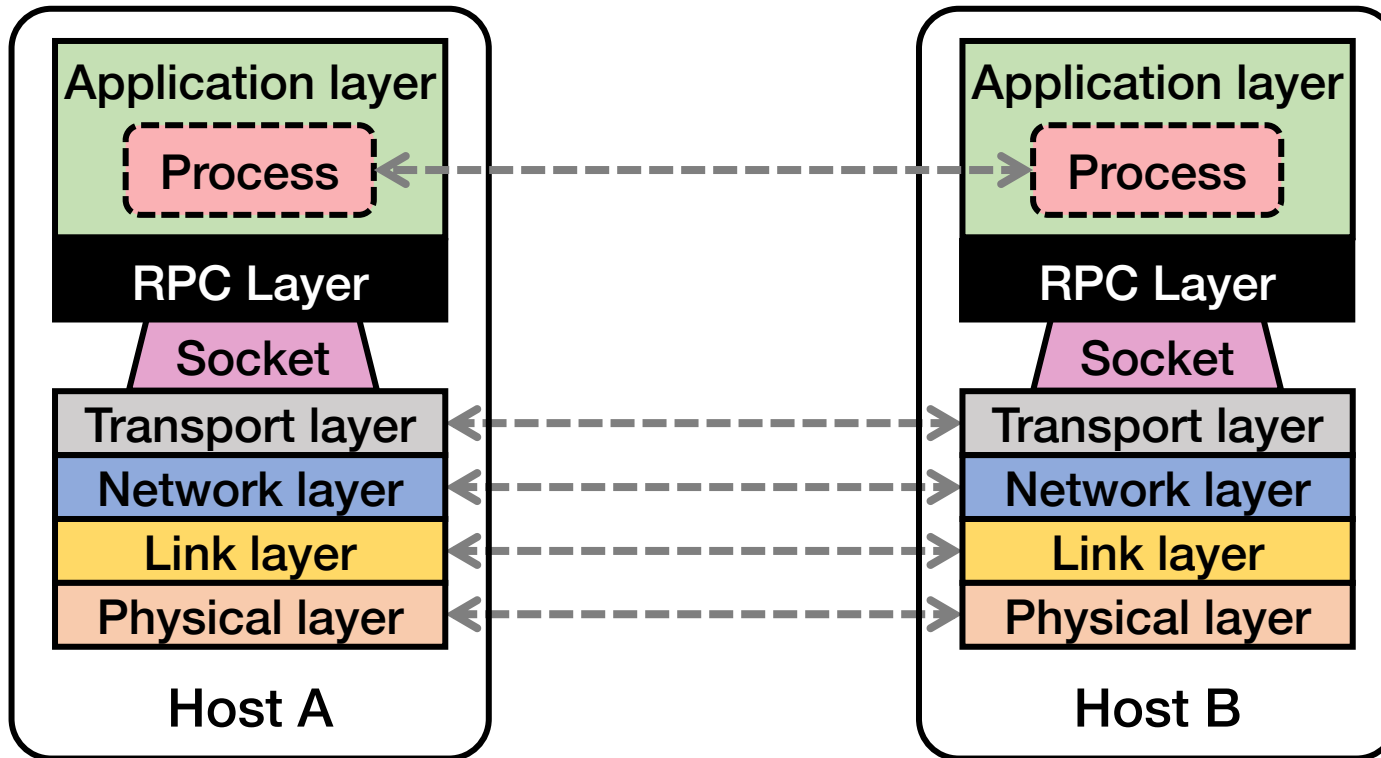
Network sockets: Summary

- Network sockets provide apps with point-to-point communication between processes
 - Only interface to the network provided by the OS
- However, the socket interface forces us to design distributed applications using a read/write interface
- `put(key, value) → message with sockets`

“All problems in computer science can be solved by another level of indirection.”


— **David Wheeler**

Solution: Another layer!



RPC Goal

- To allow programs to call **procedures** located on other machines as *if they were local*



RPC Goal: To make network communication appear like a local procedure call—transparency for procedure calls

Historical note: Seems obvious in retrospect, but RPC was only invented in the '80s.

See Birrell & Nelson, "Implementing Remote Procedure Call" ... Or Bruce Nelson, Ph.D. Thesis, Carnegie Mellon University: Remote Procedure Call, 1981

"All problems in computer science can be solved by another level of indirection"
"But that usually will create another problem."

— David Wheeler

RPC issues

1. Heterogeneity

- Dealing with differences in data representation on different machines

2. Failure

- What if messages get dropped?
- What if client, server, or network fails?

3. Performance

- Local procedure call takes ≈ 10 cycles ≈ 3 ns
- RPC in a data center takes ≈ 10 μ s (10^3 X slower)
 - In the wide area, typically 10^6 X slower

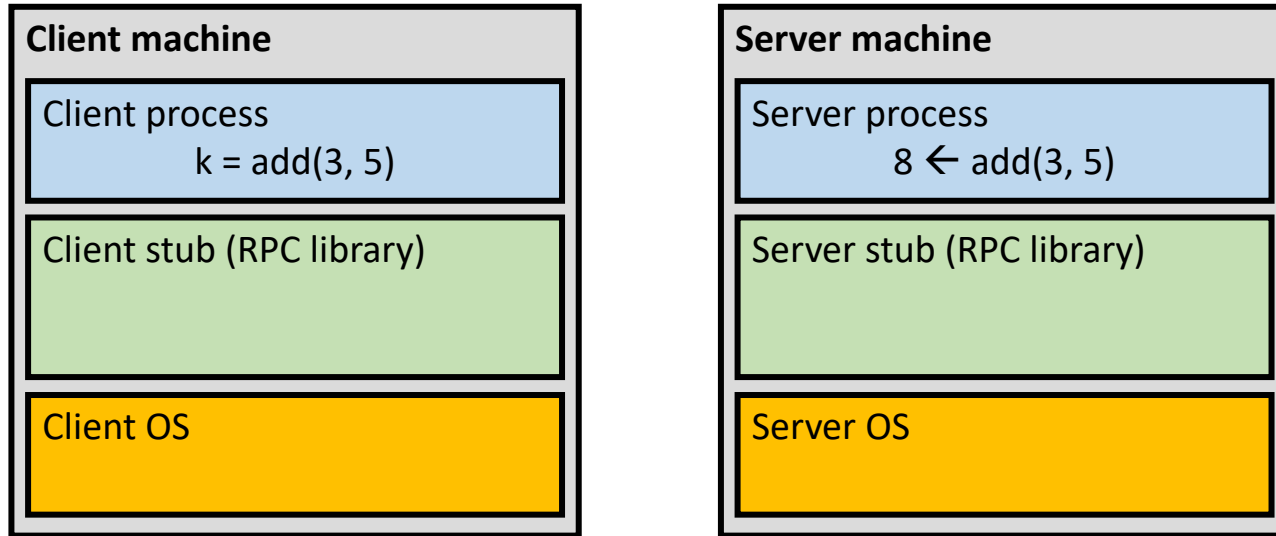
Problem: Differences in data representation

- Not an issue for local procedure calls
- For a remote procedure call, a remote machine may:
 - Use a different byte ordering (endianness)
 - Run process written in a different language
 - Represent data types using different sizes
 - Represent floating point numbers differently

Solution: Interface Description Language (IDL)

- Mechanism to pass procedure parameters and return values in a **machine-independent way**
 - E.g., Protocol Buffers, Apache Thrift
- A programmer may write an **interface description** in the IDL
 - Defines API for procedure calls: names, parameter/return types
- Then runs an **IDL compiler** which generates:
 - Code to **marshal** (convert) native data types into machine-independent byte streams
 - And vice-versa, called **unmarshaling**

How do RPCs work?

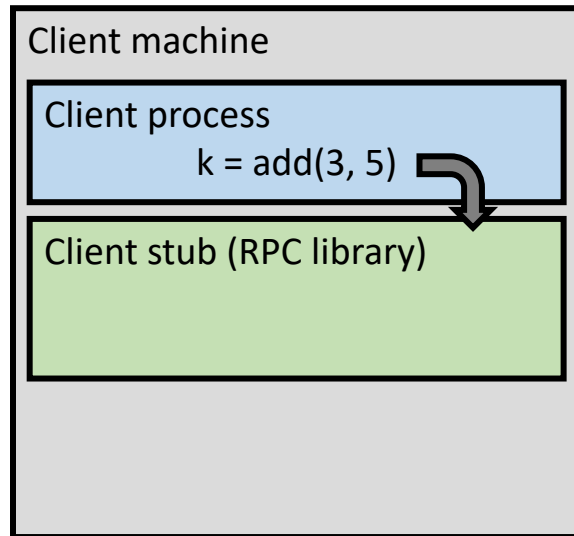


- **Client stub:** Forwards local procedure call as a request to the server
- **Server stub:** Dispatches RPC to its implementation

A day in the life of an RPC

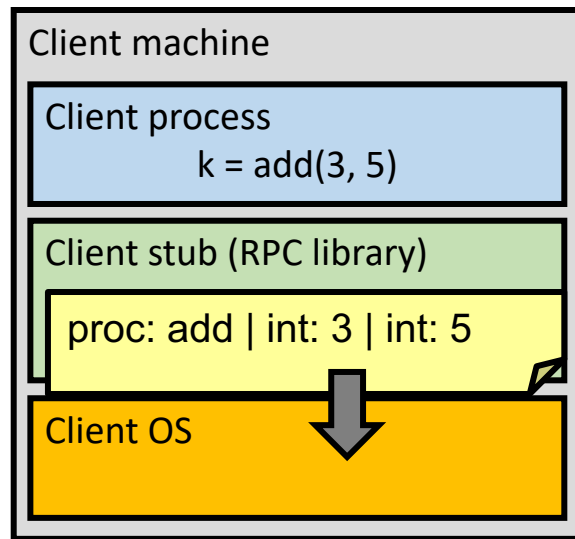
A day in the life of an RPC

1. Client calls stub function (pushes parameters onto stack)




A day in the life of an RPC

1. Client calls stub function (pushes parameters onto stack)
2. Stub **marshals** parameters to a network message



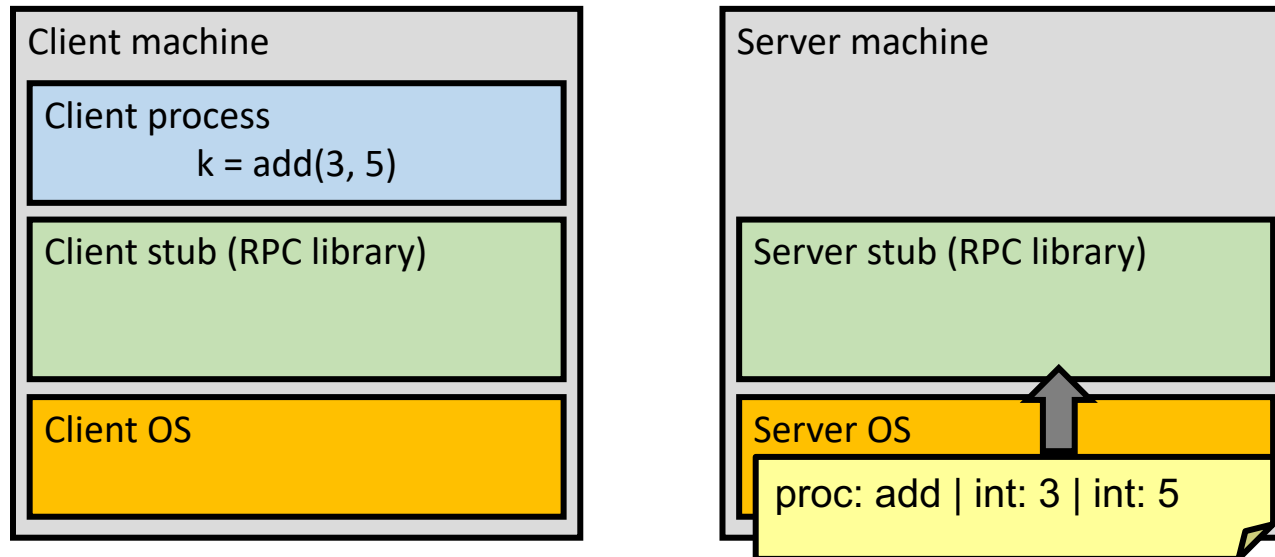
A day in the life of an RPC

2. Stub marshals parameters to a network message
3. OS sends a network message to the server



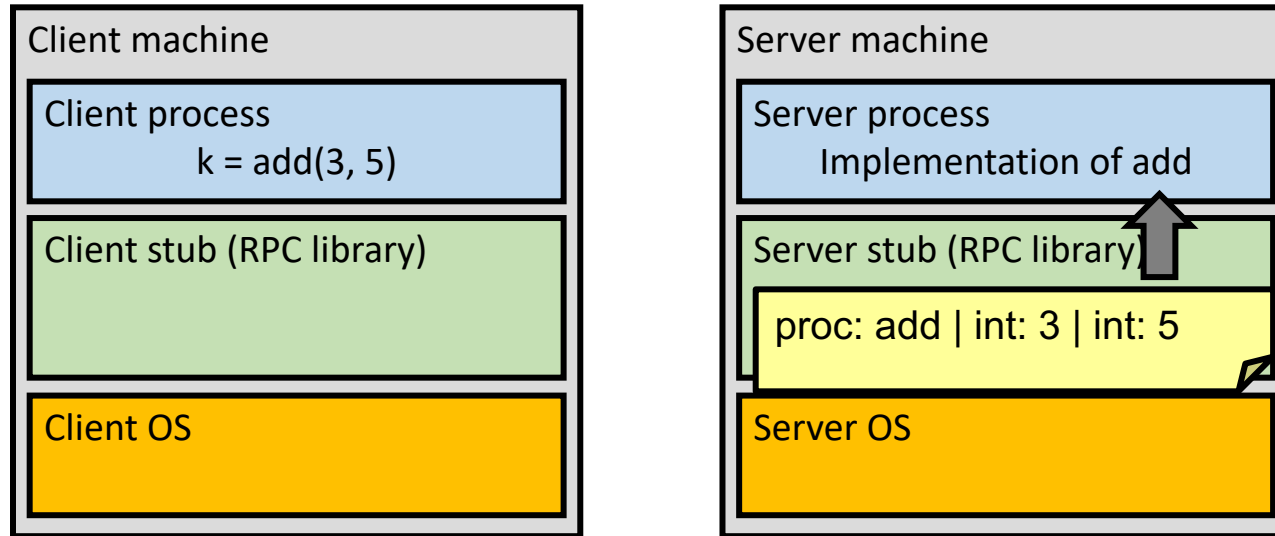
A day in the life of an RPC

3. OS sends a network message to the server
4. Server OS receives message, sends it up to stub



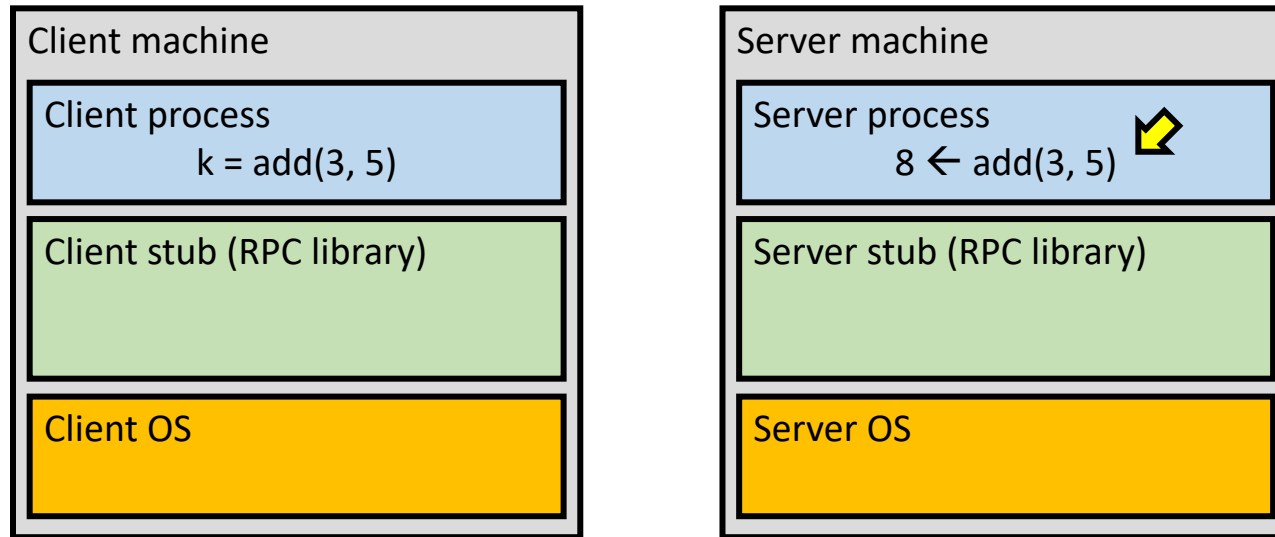
A day in the life of an RPC

4. Server OS receives message, sends it up to stub
5. Server stub unmarshals params, calls server function



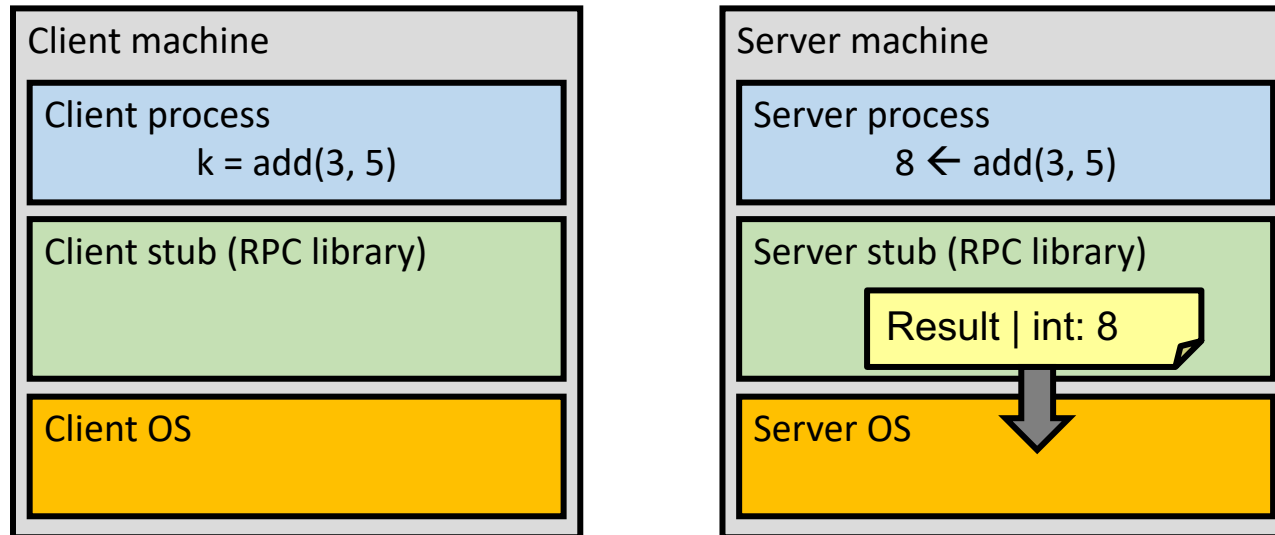
A day in the life of an RPC

5. Server stub unmarshals params, calls server function
6. Server function runs, returns a value



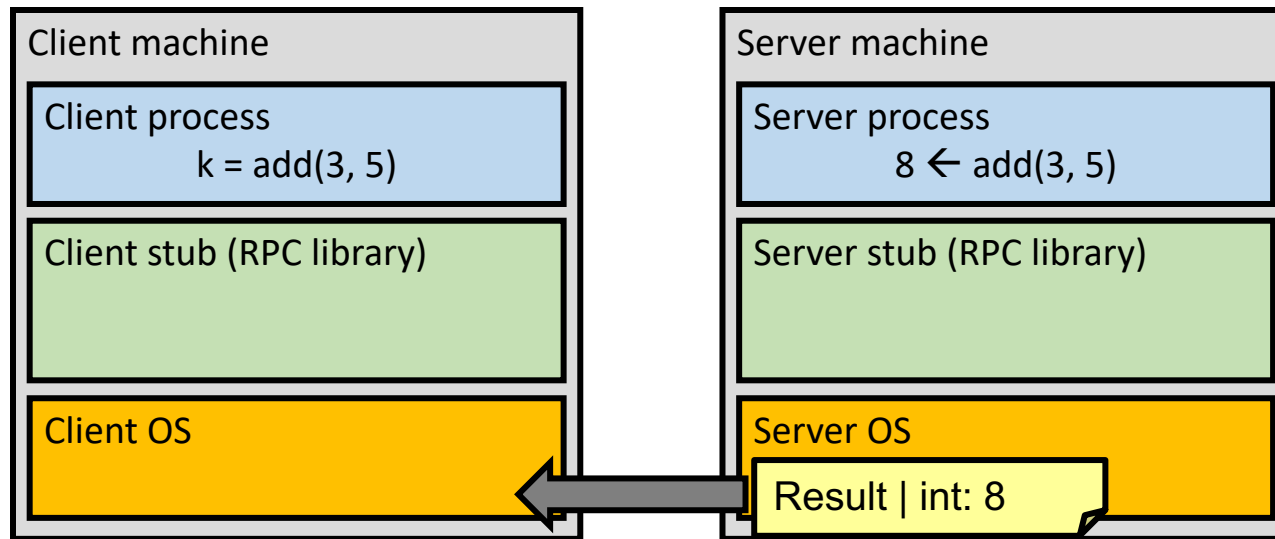
A day in the life of an RPC

6. Server function runs, returns a value
7. Server stub marshals the return value, sends message



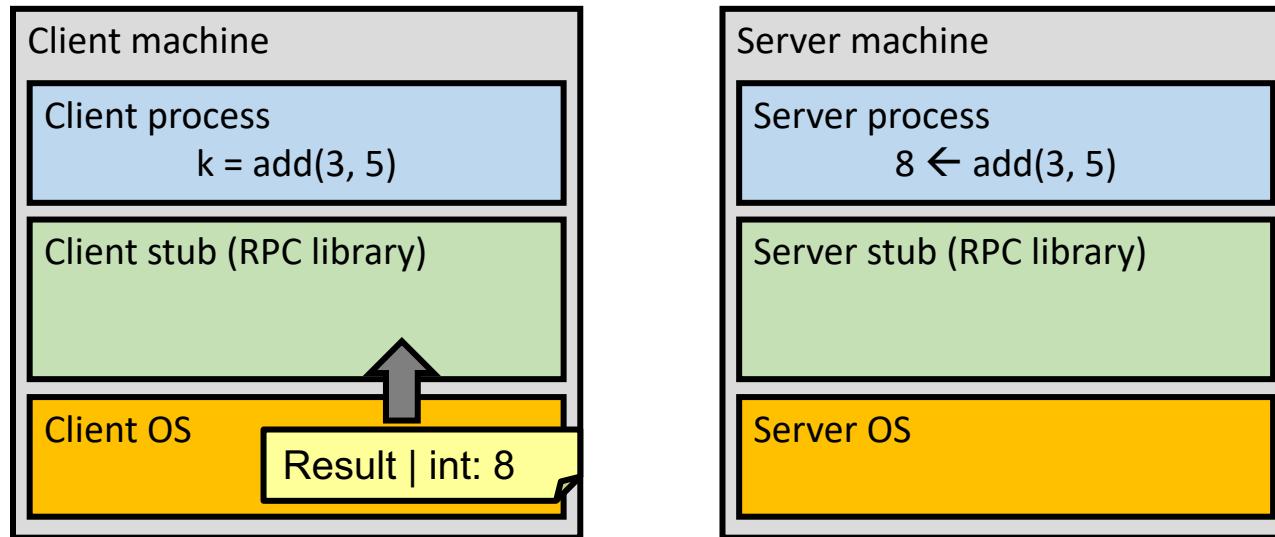
A day in the life of an RPC

7. Server stub marshals the return value, sends message
8. Server OS sends the reply back across the network



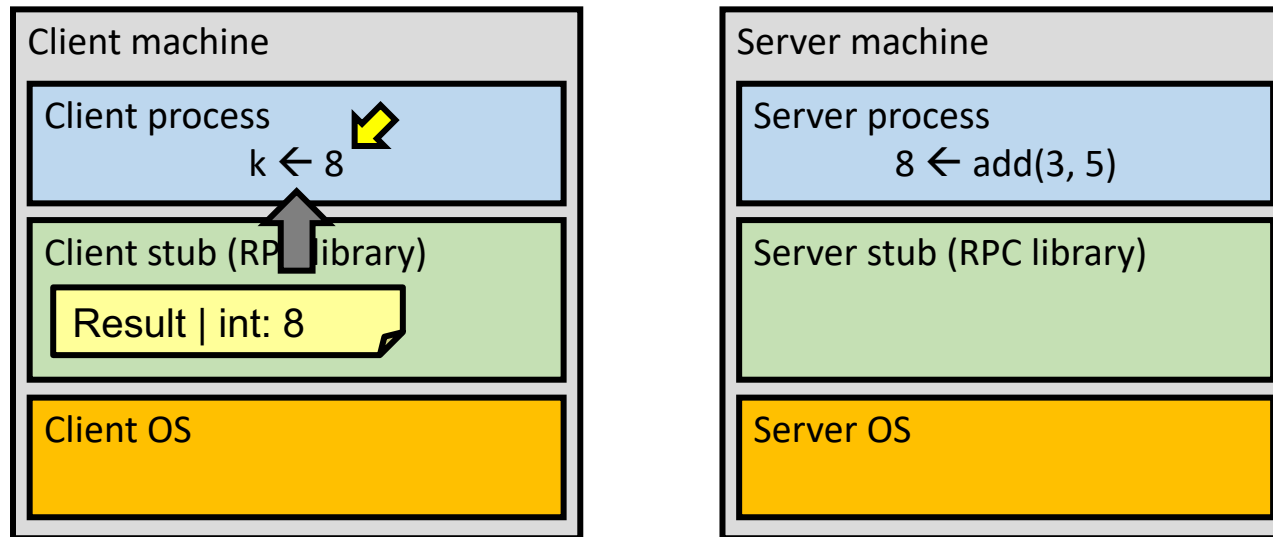
A day in the life of an RPC

8. Server OS sends the reply back across the network
9. Client OS receives the reply and passes up to stub



A day in the life of an RPC

9. Client OS receives the reply and passes up to stub
10. Client stub unmarshals return value, returns to client



Passing Value Parameters

- Copy-by-value

What about passing parameters by reference?

RPC: Passing Reference Parameters

- Replace copy-by-reference by call-by-copy/restore

Question: Would this always give us the same results as passing by reference?

What to bind to?

- Need to locate a remote host and the proper process (port)
- Solution 1:
 - Maintain a centralized database that can locate a host that provides a type of service (proposed by Birrell and Nelson in 1984)
 - A server sends a message to a central authority stating its willingness to accept certain remote procedure calls
 - Clients contact this central authority when they need to locate a service
- Solution 2:
 - Require the clients to know which hosts they need to contact

Rest of the lecture

- Failure Handling in RPCs

Failures during RPCs

What do we want in RPC systems?

- Exactly-once semantics, like local procedure calls

RPC Semantics

- Exactly-Once
 - Remote procedure will be executed exactly once on the server
- At-Least-Once
 - Remote procedure may execute more than once
 - Possible Side effects
- At-Most-Once:
 - Remote procedure will not execute more than once

At-Most-Once RPC semantics

- Idea: server RPC code detects duplicate requests
 - Returns previous reply instead of re-running the procedure

Detecting duplicate requests

- Client includes unique **transaction ID (xid)** with each RPC requests
- Client uses same xid for retransmitted requests

At-Most-Once Server

```
if seen[xid]:  
    retval = old[xid]  
else:  
    retval = remote_proc()  
    old[xid] = retval  
    seen[xid] = true  
return retval
```


At-Most-Once: Providing unique XIDs

- Combine a unique client ID with a sequence number
 - MAC address + a strictly increasing number
 - Static IP address + current time of day
- What might not work?
 - Big random number (probabilistic, no guarantees that it will be unique)
 - DHCP assigned IP + seq. number (IP might change across reboots)
 - If the client crashes and restarts, it will have a different client ID

At-Most-Once: Issues?

At-Most-Once: Discarding server state

- **Problem:** seen and old tables will grow without bound
- **Observation:** By construction, when the client gets a response to a particular xid, it will never re-send it
- **Client could tell server “I’m done with xid x – delete it”**
 - Have to tell the server about each and every retired xid
 - Could piggyback on subsequent requests

At-Most-Once: Discarding server state

- **Problem:** server state will grow without bound
- Suppose $xid = \langle \text{unique client id, sequence no.} \rangle$
 - e.g. $\langle 42, 1000 \rangle, \langle 42, 1001 \rangle, \langle 42, 1002 \rangle$
- Client includes “seen all replies $\leq X$ ” with every RPC
 - Much like TCP sequence numbers, acks

At-Most-Once: Concurrent Updates

- **Problem:** How to handle a duplicate request while the original is still executing?
 - Server doesn't know reply yet. Also, we don't want to run the procedure twice
- **Idea:** Add a `pending` flag per executing RPC
 - Server waits for the procedure to finish or ignores

At-Most-Once: Server crash and restart

- **Problem:** Server may crash and restart
- Does server need to write its state to disk (persistent memory)?
- **Yes! On server crash and restart:**
 - If state is only in volatile memory (e.g., RAM):
 - Server will forget, **accept duplicate requests**

Go's net/rpc is at-most-once

- Opens a TCP connection and writes the request
 - TCP may retransmit but server's TCP receiver will filter out duplicates internally, with sequence numbers
 - No retry in Go RPC code (i.e., will not create a second TCP connection)
- However: Go RPC returns an error if it doesn't get a reply
 - After a TCP timeout
 - Perhaps server didn't see request
 - Perhaps server processed request but server/net failed before reply came back

Exactly-Once RPC semantics?

- Need **retransmission**
- Plus the **duplicate filtering** of at most once scheme
 - To survive client crashes, client also needs to make sure it has the same unique id after restart OR
 - The client should record pending RPCs on disk
 - So it can replay them with the same unique identifier
- Plus story for making **server reliable**
 - Even if server fails, the system needs to continue with full state
 - To survive server crashes, server should log to disk results of completed RPCs (to suppress duplicates)

Synchronous and asynchronous RPCs

- Synchronous remote procedure call is a blocking call
 - I.e., when a client has made a request to the server, the client will wait until it receives a response from the server
- Asynchronous RPC: Client does not wait for a response
 - This is useful when the RPC call is a long-running computation on the server, meanwhile, client can continue execution.

Summary of RPCs

- RPCs key building block, used commonly
- Make network comm. appear like a **local procedure call**
- Issues surrounding **machine heterogeneity**
- Subtle issues around **failures**
 - Need retransmissions to deal with failures
 - At-most-once w/ duplicate filtering
 - Discard server state w/ cumulative acks
 - Exactly-once with:
 - retransmissions + at-most-once fault tolerance (of servers)

