


CS 582: Distributed Systems

Vector Clocks & Leader Elections



Dr. Zafar Ayyub Qazi
Fall 2024

This lecture will have a few slippery slopes



Today's Agenda

- Application of Vector Clocks
 - Causally-ordered Slack-like application
- Leader Elections
- Leader Elections in Raft
 - Quick overview of Raft
 - Leader election algorithm

Specific learning outcomes

By the end of today's lecture, you should be able to:

- ☐ Apply Vector clocks to provide causally ordered communication
- ☐ Describe the leader election problem
- ☐ Explain how leader elections are conducted in Raft
- ☐ Analyze leader election algorithm in Raft for safety and liveness

Recap: Motivation for Vector Clocks

- How can we design logical clocks that can allow us to infer potential causality?
 - More precisely, whether $a \rightarrow b$?

Recap: Vector Clocks

- A **Vector Clock** is a **vector of integers**, one entry for each process in the system
- Label each event e with a vector $V(e) = \langle c_1, c_2, \dots, c_n \rangle$
 - Where c_i is the count of events in process i that “happen-before” e
 - And we have n processes in the system
- **Initially**, all vectors are $\langle 0, 0, \dots, 0 \rangle$

Recap: Rules for updating Vector Clocks

- Local update rule:
 - For each local event on process i , increment local entry c_i (by 1)
- Message rule:
 - If process j receives message with vector $\langle d_1, d_2, \dots, d_n \rangle$:
 - Set each local entry $c_k = \max\{c_k, d_k\}$
 - Increment local entry c_j (by 1)

Vector Clock Orderings


- Rules for comparing vector timestamps:
 - $V(a) = V(b)$ iff $a_k = b_k$ for all k
 - $V(a) < V(b)$ iff $a_k \leq b_k$ for all k and $V(a) \neq V(b)$
 - $V(a) \parallel V(b)$ iff $a_i < b_i$ and $a_j > b_j$, for some i, j
- Properties of this order
 - $(V(a) < V(b)) \iff (a \rightarrow b)$
 - $(V(a) = V(b)) \iff (a = b)$
 - $(V(a) \parallel V(b)) \iff (a \parallel b)$

Vector Clocks Capture Potential Causality

- $V(w) < V(z)$ then there is a chain of events linked by happens-before (\rightarrow) between w and z

Vector Clocks Capture Potential Causality

- $V(w) < V(z)$ then there is a chain of events linked by happens-before (\rightarrow) between w and z
- $V(a) \parallel V(w)$ then there is **no** such chain of events between a and w



Two events a, z


Lamport Clocks: $C(a) < C(z)$

Conclusion: **None** $\overline{a \rightarrow b \text{ or } a \parallel z}$

Vector Clocks: $V(a) < V(z)$

Conclusion: $a \rightarrow \dots \rightarrow z$


Vector clock timestamps precisely capture the happens-before relation



Slack-Like Application


- Slack-like application
 - Broadcast posts to all other users
- **Goal:** No user should display a response post before the corresponding original message post
- Deliver a message only after all messages that might causally precede it have been delivered
 - Otherwise, the user would see a reply to a message they could not find

Causally-ordered Slack-like App



- User 1 posts
- User 2 replies to User 1's post
- User 3 observes

Causally-ordered Slack-like App



For enforcing causal message delivery, we assume clocks are adjusted only when sending and **delivering** messages **delivering** messages to application layer


1. Upon sending a message, process P_i will increment $V_i[i]$ by 1
2. When it delivers a message m with timestamp $ts(m)$, it only adjusts $V_i[k]$ to $\max\{V_i[k], ts(m)[k]\}$ for each k

- User 1 posts
- User 2 replies to User 1's post
- User 3 observes

Physical time
Suppose that P_j receives a message m from P_i . Delivery of the message to the application layer of a process P_j is delayed until the following two conditions are met:

1. $ts(m)[i] = V_j[i] + 1$
2. $ts(m)[k] \leq V_j[k]$ for all $k \neq i$

Causally-ordered Slack-like App



For enforcing causal message delivery, we assume clocks are adjusted only when sending and **delivering** messages to application layer

1. Upon sending a message, process P_i will increment $V_i[i]$ by 1
2. When it delivers a message m with timestamp $ts(m)$, it only adjusts $V_i[k]$ to $\max\{V_i[k], ts(m)[k]\}$ for each k

- User 1 posts
- User 2 replies to User 1's post
- User 3 observes

Suppose that P_j receives a message m from P_i . Delivery of the message to the application layer of a process P_j is delayed until the following two conditions are met:

1. $ts(m)[i] = V_j[i] + 1$
2. $ts(m)[k] \leq V_j[k]$ for all $k \neq i$

Leader Elections



Why Election?

- Often need one process to play a special role
 - Coordinator, Initiator, etc.
 - Useful for coordination among distributed servers
 - E.g., pick a server node in Berkeley algorithm
- How to select a leader if you have multiple candidates?
- Leader Election widely used in industry
 - Apache Zookeeper, Google's Chubby, Google Spanner database, Network Time Protocol, Berkeley Algorithm, Raft, etc.

Leader Election Problem

- In a group of processes, elect a **Leader** to undertake special tasks
 - And let **everyone** in the group know about this Leader
- What happens when a leader fails (crashes)?
 - Some process detects this (e.g., using timeouts)
 - Then what?
- Need to start elections to elect a leader
 - An election algorithm defines how elections are conducted

Goals of an Election Algorithm

1. Elect only one leader among the non-faulty processes
2. All non-faulty processes agree on who the leader is

Requirements

- Any process can call for an election
- A process can call for at most one election at a time
- Multiple processes are allowed to call an election simultaneously
 - All of them together must yield only a single leader

Leader Elections in Raft

Short overview of Raft

- We will come back and study Raft in detail a couple of weeks later

What is Raft?

- Raft is a distributed consensus algorithm
- Consensus algorithms allow a collection of machines to work as a coherent group that can survive the failures of some of its members
 - Play a key role in building large-scale distributed software systems
- Raft is deployed by several companies
- You will implement it in assignments 2-4

Replicated State Machines

- Consensus algorithms typically arise in the context of replicated state machines

Replicated State Machines



Clients






Servers

Replicated State Machines

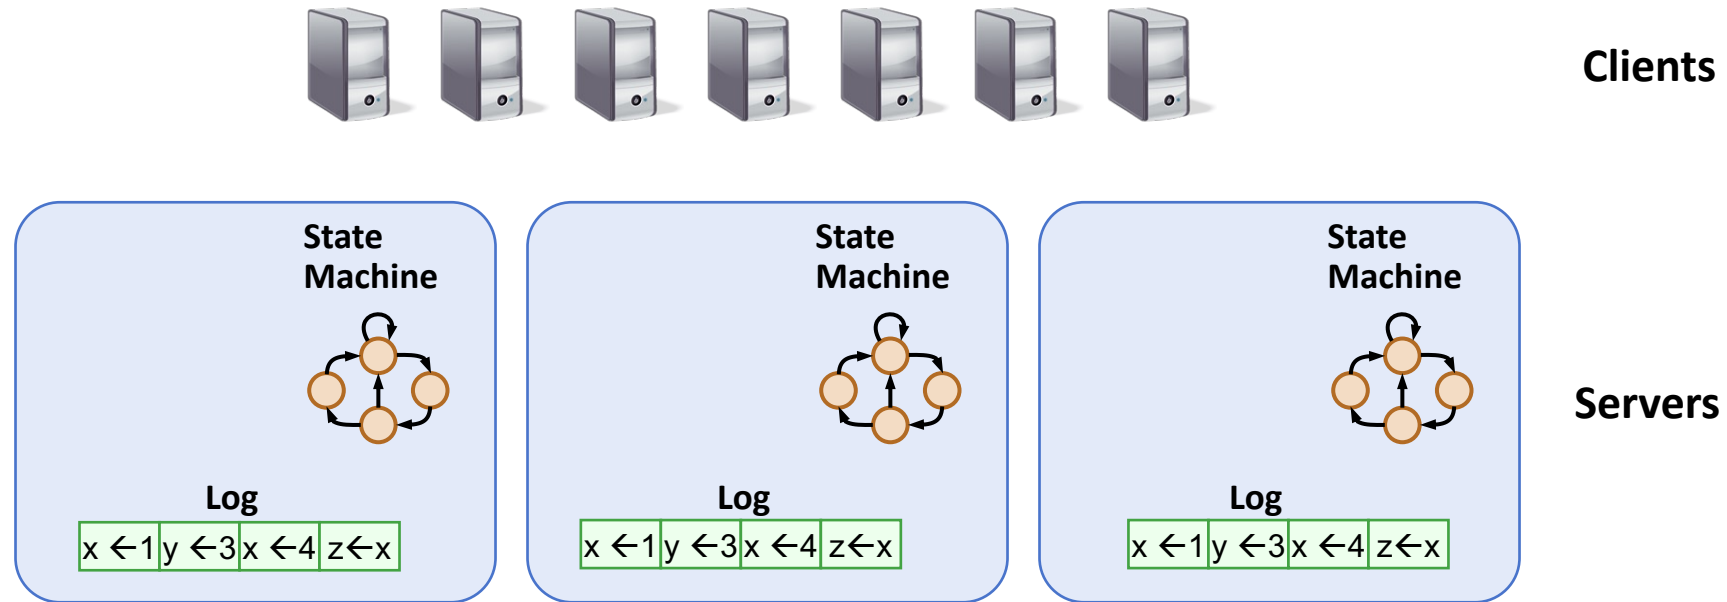


Clients



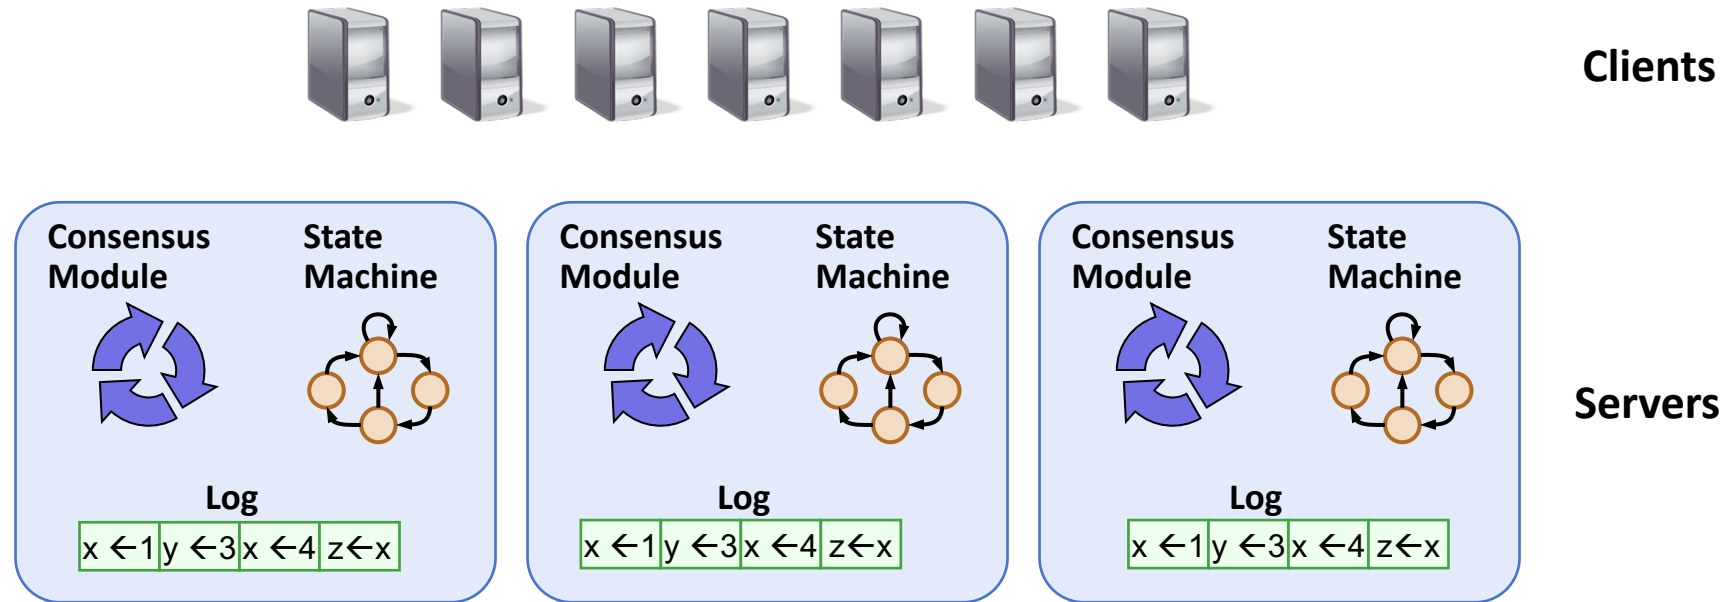
Servers

Raft → Replicated Log



- Replicated log => replicated state machine
 - All servers execute same commands in same order

Raft → Replicated Log



- Replicated log => replicated state machine
 - All servers execute same commands in same order
- Consensus module ensures proper log replication
- System should make progress as long as any majority of servers are up
- Failure model: fail-stop (not Byzantine), delayed/lost messages

Raft has multiple components

- Leader Election
- Log Replication
- Safety
- Membership change

Raft servers can be in three states

Follower

Only responds to incoming requests

Candidate

Used to elect a new leader

Leader

Handles all client interactions, log replication

Servers use RPCs

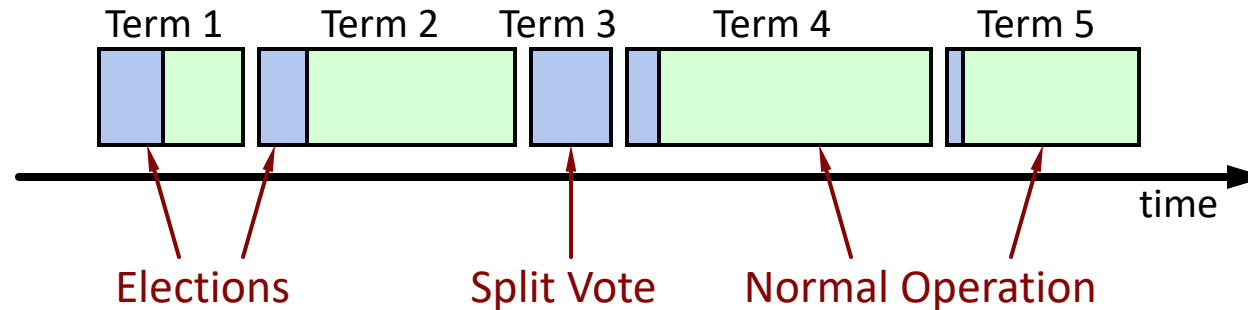
- RequestVote RPCs
 - Sent when a candidate is requesting votes to become a leader
- AppendEntries RPCs
 - To send heartbeat messages: to let other servers know “I am up”
 - Another purpose is log replication, which we will not discuss today

Time is divided into Terms





Time is divided into Terms




- Time divided into terms:
 - Election
 - Normal operation under a single leader
- At most 1 leader per term
- Some terms have no leader (failed election)
- Each server maintains current term value (no global view)
 - Exchanged in every RPC
 - Peer has later term? Update term, revert to follower
 - Incoming RPC has obsolete term? Reply with error

How are elections conducted in Raft?

Election Basics


- If a server suspects a leader has failed, **increment current term**
- Change to Candidate state
- Vote for self
- Send **RequestVote** RPCs to all other servers, retry until either:
 1. **Receive votes from majority of servers:**
 - Become leader
 - Send AppendEntries heartbeats to all other servers
 2. **Receive RPC from a valid leader:**
 - Return to follower state
 3. **No-one wins election (election timeout elapses):**
 - Increment term, start new election

Server State Transitions



Election Correctness

- **Safety:** allow at most one winner per term
 - Each server gives out only one vote per term (persist on disk)
 - Two different candidates can't accumulate majorities in same term



- **Liveness:** some candidate must eventually win
 - Choose election timeouts randomly in $[T, 2T]$ (e.g., 150-300 ms)
 - One server usually times out and wins election before others wake up
 - Works well if $T \gg$ broadcast time

Summary: Leader Election in Raft

- Three possible server states: follower, candidate, leader
- All servers startup as followers
- If a follower detects a leader has failed
 - Follower can start new election
 - Increments term, and transition to Candidate state
 - Candidate becomes a leader if a majority of processes vote for it
 - Possible no one wins election: split votes

For more on leader elections in Raft ...

- Attend the [next tutorial](#)
- Check out [Raft's extended paper](#) (will be shared with the assignment 2 handout)
- Checkout the following [excellent visualization of elections in Raft](#):
 - <http://thesecretlivesofdata.com/raft/>

Next Lecture

- Replication and Consistency
 - Consistency Models
 - CAP