

CS 582: Distributed Systems

Paxos



Dr. Zafar Ayyub Qazi

Fall 2024

A note before we start

- This is the most important module of the course and the most tricky
 - We will cover several consensus protocols, e.g., Paxos, Raft, PBFT,
- An important goal: Build strong intuition for the key ideas
 - What are the key ideas? Why do they work? When they might not work?
- Requirement from you:
 - Attend classes, pay attention, and ask questions
 - Do recommended readings, homework questions, and assignment
 - Review these concepts after class to consolidate your understanding

Today's Agenda

- Basic Paxos

Specific learning outcomes

By the end of today's lecture, you should be able to:

- ☐ Explain how Basic Paxos works
- ☐ Analyze the design choices made in Basic Paxos
- ☐ Analyze and evaluate Paxos in terms of safety, liveness, and fault tolerance

Recap: Consistency Protocols


- Single phase primary-backup schemes
 - Safety: Can't roll back updates, so safety can get violated
 - Liveness: doesn't provide liveness with failures
- Two Phase Commit (2-PC)
 - Safety: Allows for rolling back updates and making nodes consistent
 - Liveness: blocks in case of failures

Paxos

Paxos

- Paxos
 - Consensus algorithm
 - Proposed by Leslie Lamport in 1989
- Paxos and its variants widely used in industry
 - Zookeeper (Yahoo), Google Chubby, Google Spanner and many others

Raft Goal → Replicated Log



- Replicated log => replicated state machine
 - All servers execute same commands in same order
- Consensus module ensures proper log replication
- System should make progress as long as any majority of servers are up
- Failure model: fail-stop (not Byzantine), delayed/lost messages

The Paxos Approach

- Decompose the problem:
- Basic Paxos ("single decree"):
 - One or more servers propose values
 - System must agree on a single value as chosen
 - Only one value is ever chosen
- Multi-Paxos:
 - Uses multiple instances of basic Paxos (one for each log entry)

Requirements for Basic Paxos

- Safety:
 - Only a single value may be chosen
 - A server never learns that a value has been chosen unless it really has been
 - Only a value that has been proposed may be chosen
- Liveness (as long as majority of servers up and communicating with reasonable timeliness):
 - Some proposed value is eventually chosen
 - If a value is chosen, servers eventually learn about it

Paxos Components

- Proposers

- Active: put forth particular values to be chosen
- Handle client requests

- Acceptors

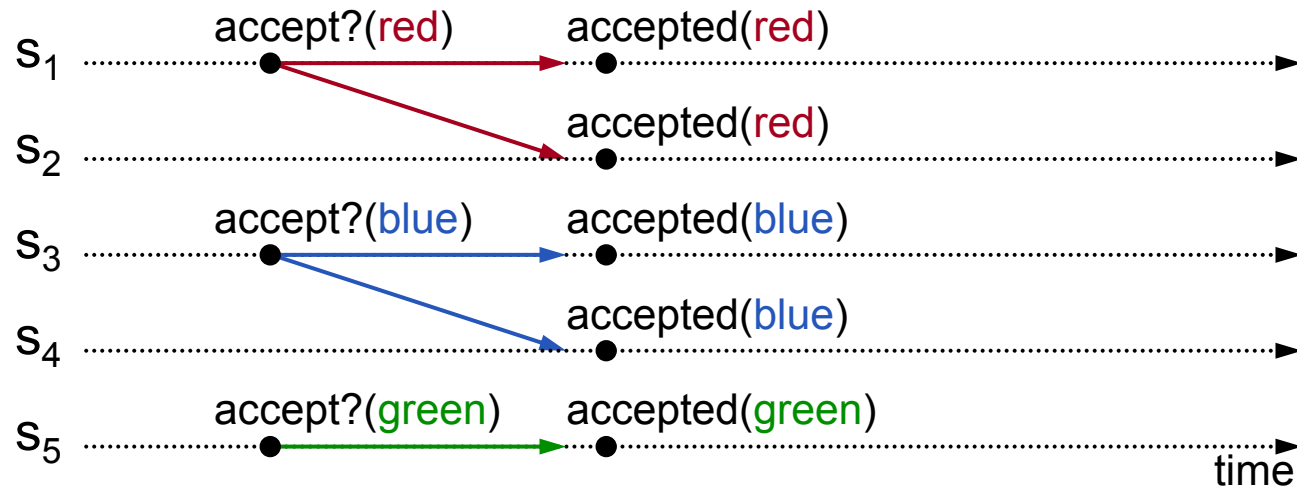
- Passive: respond to messages from proposers
- Responses represent votes that form a consensus
- Store chosen value, state of the decision process
- Want to know which value was chosen

- We will assume each Paxos server contains both components

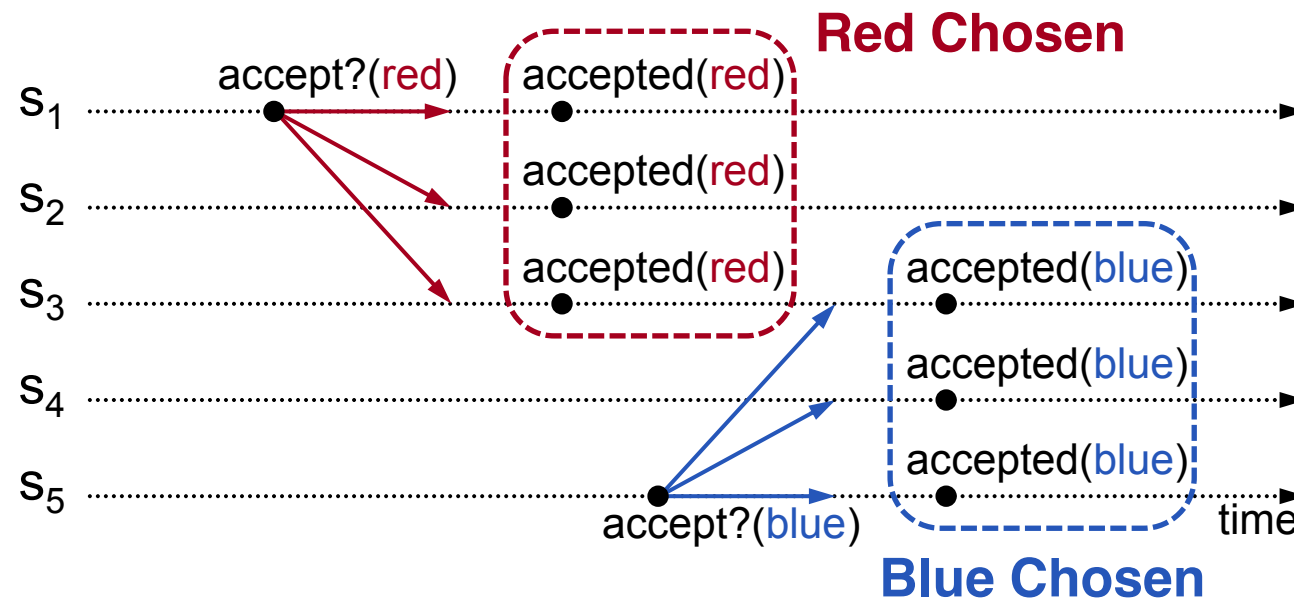
Design Decision

**How does an Acceptor accept values
(i.e., vote for a Proposer)?**

Option#1: Acceptor accepts first value it receives

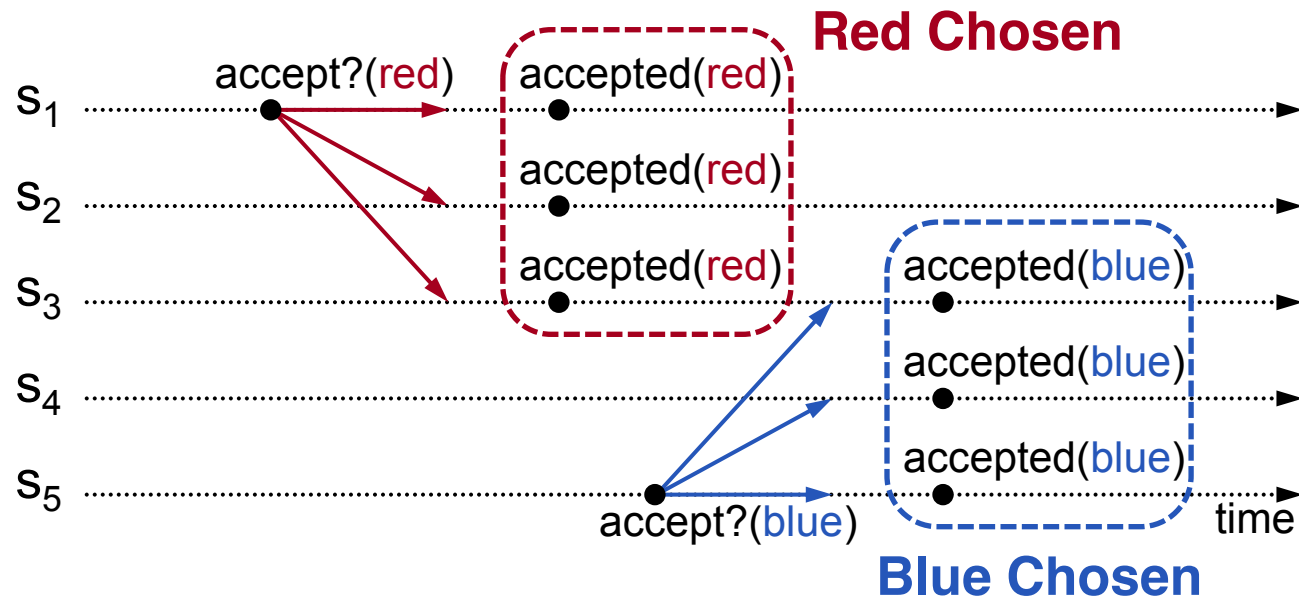


Option#2: Acceptor accepts every value it receives



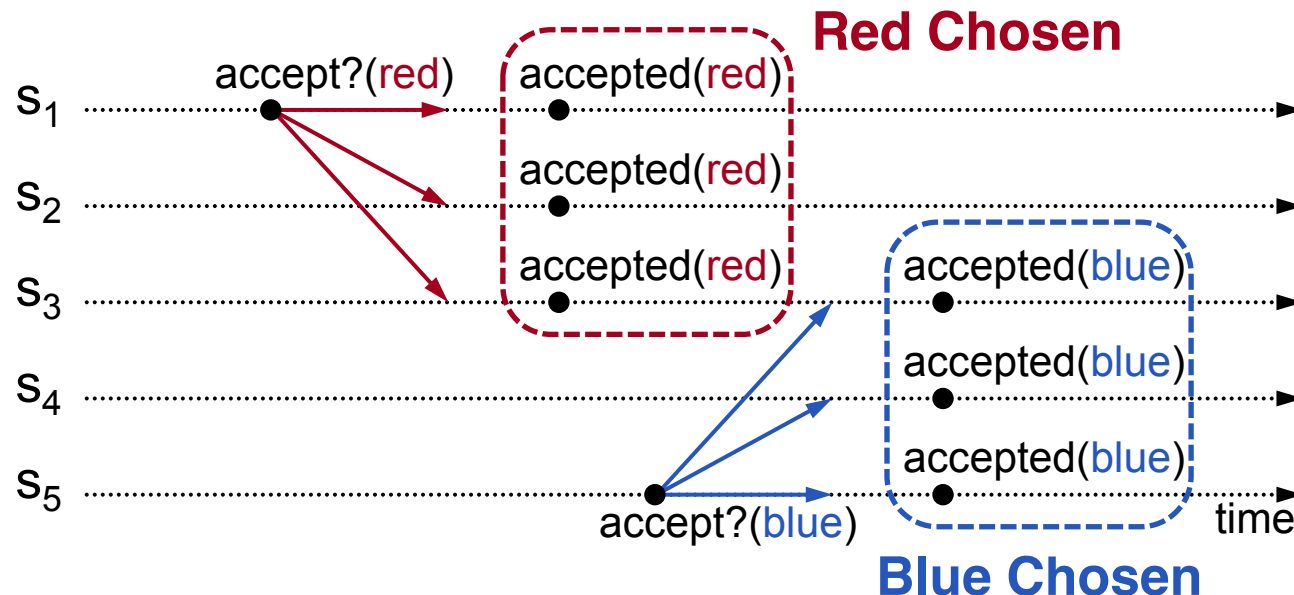
Challenge:
How to check if a value has already been chosen?

Option#1: Wait to hear from some majority of Acceptors who have accepted a value

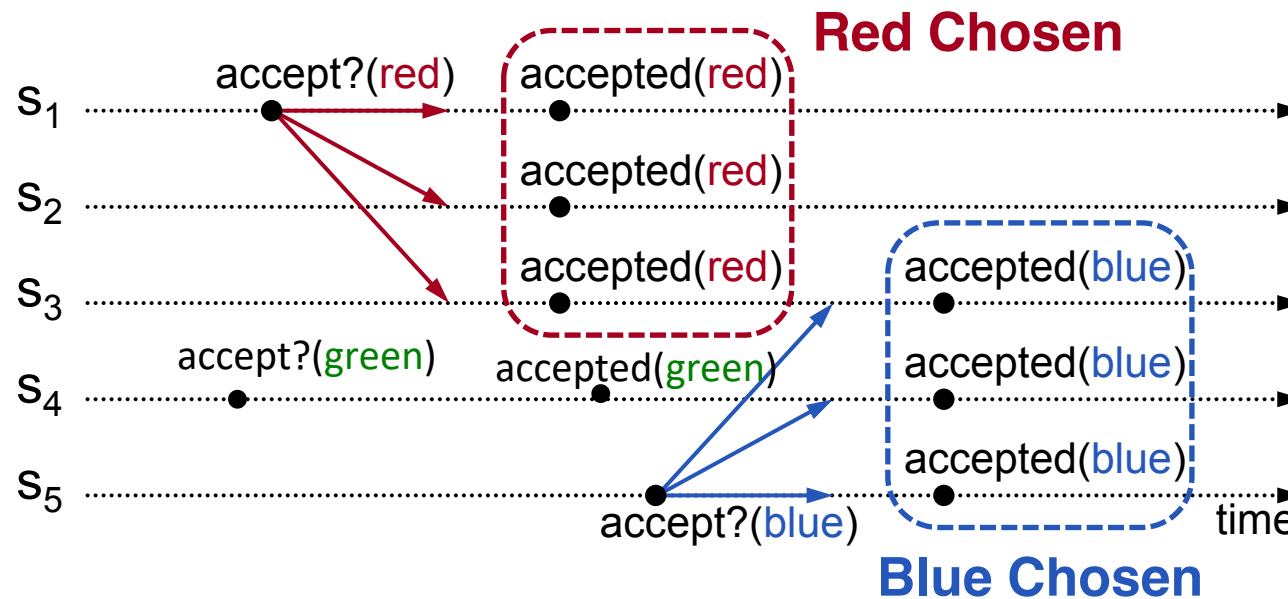


Option#2: Wait to hear from any majority of Acceptors

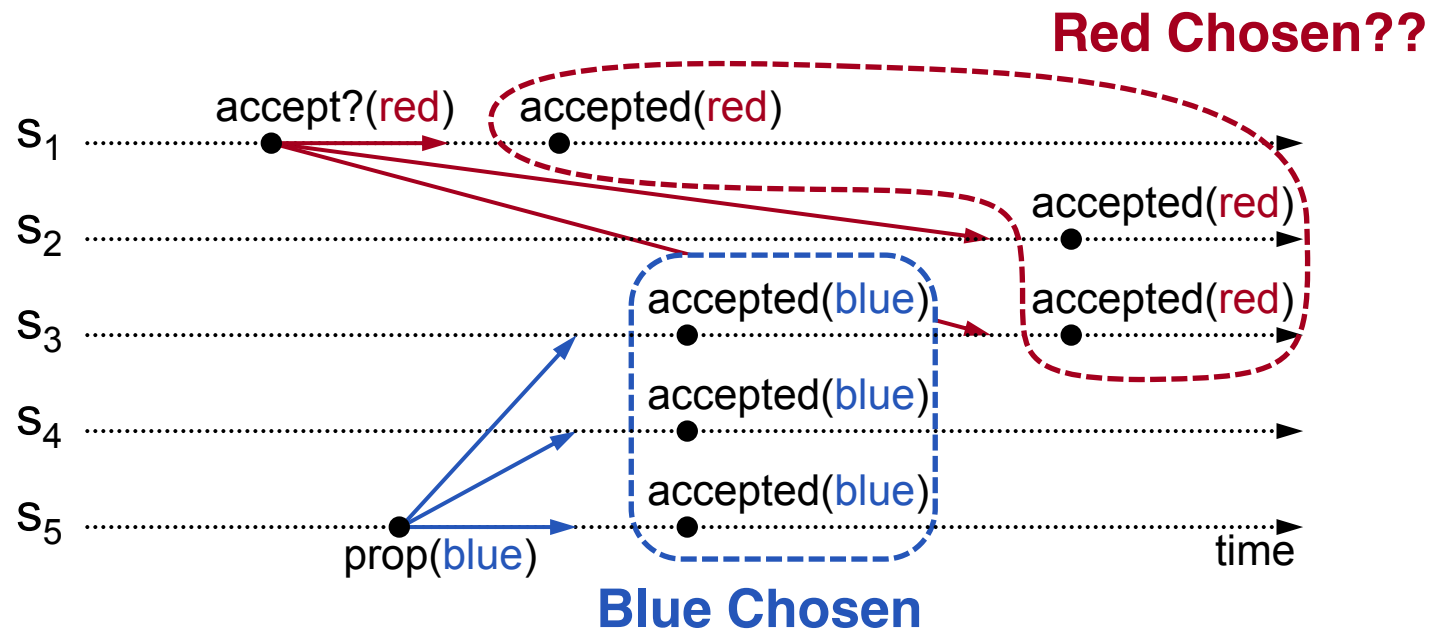
- Sub-option#1: If a majority of acceptors return an accepted value, then do not propose your value, otherwise go ahead with your proposed value
- Sub-option#2: If any of the acceptors returns an accepted value, do not propose your value, otherwise go ahead with your proposed value



Problem: What if the Proposer hears two different accepted values?



Another Problem: What if the second Proposer doesn't hear any accepts about a previous proposal?



Summary: Paxos Key Ideas

- A value is chosen only if a majority of acceptors accept it
- Use **two phases** (like in 2-PC)
 - First, check if there is a proposed value that has been chosen
 - Wait for some majority, if any acceptor has accepted a value, assume the value is chosen, and do not propose a different value
 - But two phases may not be enough
- **Order proposals and block older proposals**
 - To avoid choosing two different values and violating safety

Proposal Numbers

- Each proposal has a **unique** number
 - Higher numbers take priority over lower numbers
 - It must be possible for a proposer to choose a new proposal number higher than anything it has seen/used before
- **One simple approach:**

Proposal Number	
Round Number	Server Id

 - Each server stores **maxRound**: the largest Round Number it has seen so far
 - To generate a new proposal number:
 - Increment maxRound
 - Concatenate with Server Id
 - Proposers must persist maxRound on disk; so a proposer does not reuse proposal numbers after crash/restart

Two-Phase Approach

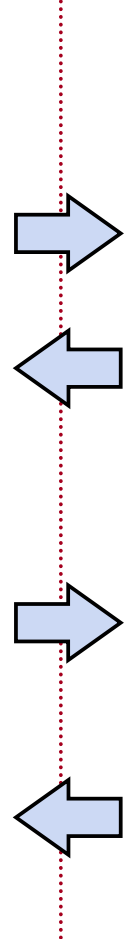
- Phase 1: Broadcast **Prepare** Message
 - Find out about any chosen values
 - Block older proposals that have not yet been completed
- Phase 2: Broadcast **Accept** Message
 - Ask acceptors to accept a specific value

Paxos

Proposers

- 1) Choose new proposal number n
- 2) Broadcast `Prepare(n)` to all servers
- 4) When responses received from majority:
 - If any `acceptedValues` returned, replace value with `acceptedValue` for highest `acceptedProposal`
- 5) Broadcast `Accept(n , value)` to all servers
- 6) When responses received from majority:
 - Any rejections (`result > n`)? goto (1)
 - Otherwise, **value is chosen**

Acceptors

- 
- 3) Respond to `Prepare(n)`:
 - If $n > \text{minProposal}$ then $\text{minProposal} = n$
 - `Return(acceptedProposal, acceptedValue)`
 - 6) Respond to `Accept(n , value)`:
 - If $n \geq \text{minProposal}$ then
 $\text{acceptedProposal} = \text{minProposal} = n$
 $\text{acceptedValue} = \text{value}$
 - `Return(minProposal)`

Acceptors must record `minProposal`, `acceptedProposal`, and `acceptedValue` on stable storage (disk)

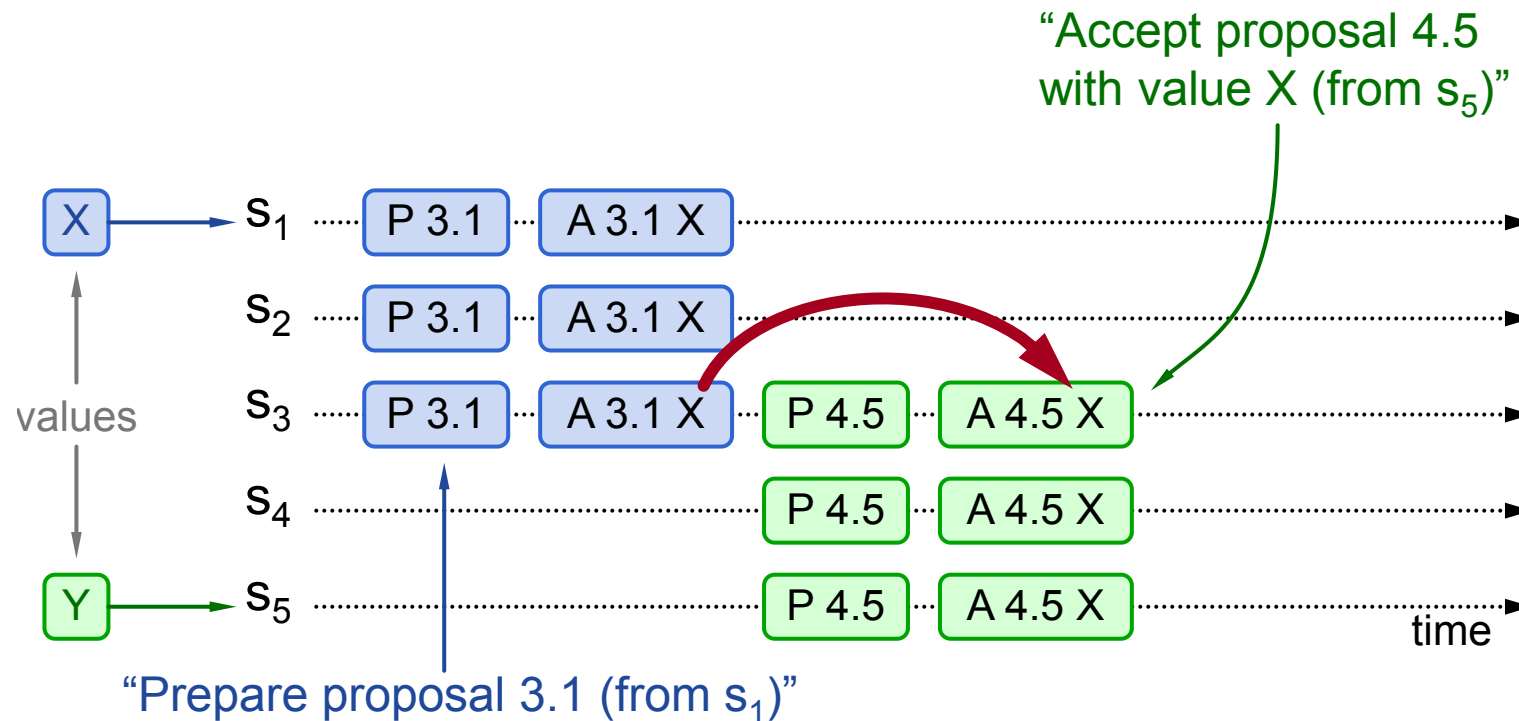
Paxos Examples

- To illustrate how Paxos works
- In particular, we will consider scenarios with two competing proposals
- There are three possible situations in such scenarios

Paxos Examples

1. Previous value already chosen

- New proposer will find it and use it



Paxos Protocol

Proposers

- 1) Choose new proposal number n
- 2) Broadcast `Prepare(n)` to all servers
- 4) When responses received from majority:
 - If any `acceptedValues` returned, replace value with `acceptedValue` for highest `acceptedProposal`
- 5) Broadcast `Accept(n , value)` to all servers
- 6) When responses received from majority:
 - Any rejections (result > n)? goto (1)
 - Otherwise, **value is chosen**

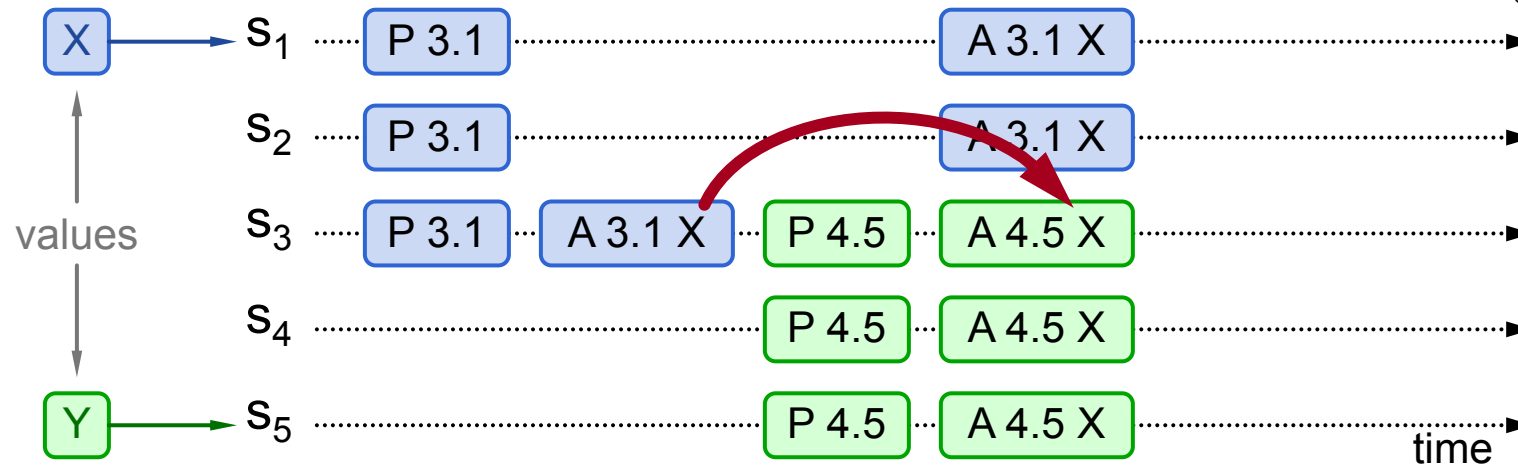
Acceptors

- 3) Respond to `Prepare(n)`:
 - If $n > \text{minProposal}$ then $\text{minProposal} = n$
 - `Return(acceptedProposal, acceptedValue)`
- 6) Respond to `Accept(n , value)`:
 - If $n \geq \text{minProposal}$ then $\text{acceptedProposal} = \text{minProposal} = n$
 $\text{acceptedValue} = \text{value}$
 - `Return(minProposal)`

Paxos Examples

2. Previous value not chosen, but new proposer sees it

- New proposer will use existing value
- Both proposers can succeed



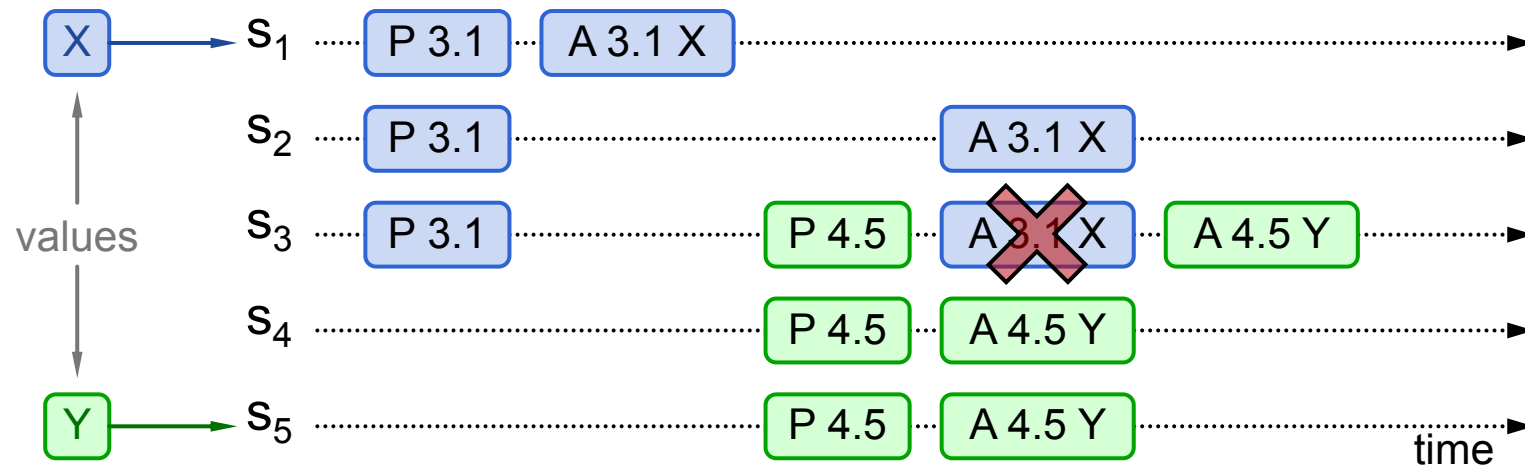
Paxos Protocol

- | Proposers | Acceptors |
|--|--|
| 1) Choose new proposal number n | |
| 2) Broadcast Prepare(n) to all servers | 3) Respond to Prepare(n) : <ul style="list-style-type: none">▪ If $n > \text{minProposal}$ then $\text{minProposal} = n$▪ Return(acceptedProposal, acceptedValue) |
| 4) When responses received from majority: <ul style="list-style-type: none">▪ If any acceptedValues returned, replace value with acceptedValue for highest acceptedProposal | |
| 5) Broadcast Accept(n, value) to all servers | 6) Respond to Accept(n, value) : <ul style="list-style-type: none">▪ If $n \geq \text{minProposal}$ then $\text{acceptedProposal} = \text{minProposal} = n$
$\text{acceptedValue} = \text{value}$▪ Return(minProposal) |
| 6) When responses received from majority: <ul style="list-style-type: none">▪ Any rejections ($\text{result} > n$)? goto (1)▪ Otherwise, value is chosen | |

Paxos Examples

3. Previous value not chosen, new proposer doesn't see it

- New proposer chooses its own value
- Older proposal blocked



Paxos Protocol

- | Proposers | Acceptors |
|--|---|
| 1) Choose new proposal number n | |
| 2) Broadcast <code>Prepare(n)</code> to all servers | 3) Respond to <code>Prepare(n)</code> : <ul style="list-style-type: none">▪ If $n > \text{minProposal}$ then $\text{minProposal} = n$▪ <code>Return(acceptedProposal, acceptedValue)</code> |
| 4) When responses received from majority: <ul style="list-style-type: none">▪ If any <code>acceptedValues</code> returned, replace value with <code>acceptedValue</code> for highest <code>acceptedProposal</code> | 6) Respond to <code>Accept(n, value)</code> : <ul style="list-style-type: none">▪ If $n \geq \text{minProposal}$ then $\text{acceptedProposal} = \text{minProposal} = n$
$\text{acceptedValue} = \text{value}$▪ <code>Return(minProposal)</code> |
| 5) Broadcast <code>Accept(n, value)</code> to all servers | |
| 6) When responses received from majority: <ul style="list-style-type: none">▪ Any rejections ($\text{result} > n$)? goto (1)▪ Otherwise, value is chosen | |

Other points to note

- Only proposer knows which value has been chosen
- If other servers want to know, must execute Paxos with their own proposal

Paxos

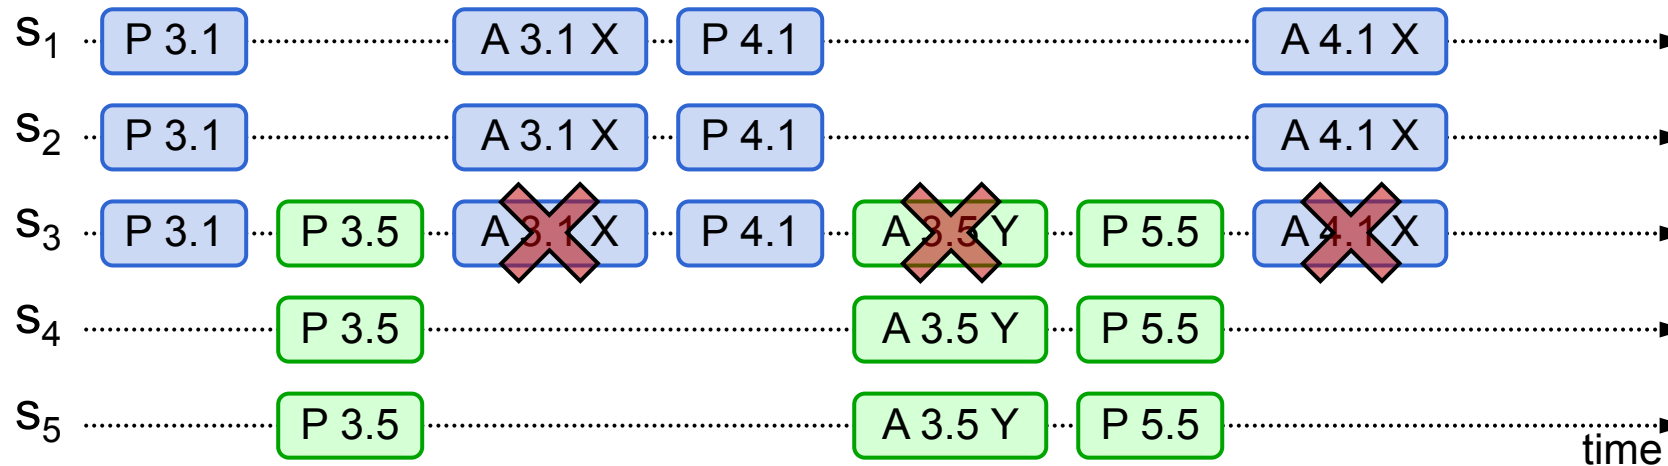
- Safety?
- Liveness?
- Performance?

Safety

- **Intuition:** if a proposal with value v is chosen, then every higher-numbered proposal issued by any proposer has value v

Liveness

- Competing proposers can livelock:



- One solution: **randomized delay** before restarting
 - Give other proposers a chance to finish choosing
- Can use **leader elections**

Performance?

Paxos Fault tolerance

- If there can be f fail-stop failures in a system
- What are the minimum number of nodes Paxos needs to ensure consensus is reached?

Paxos: Summary

- Safety: Never violated
- On Liveness
 - If things go well sometime in the future (messages and failures, etc.), there is a good chance consensus will be reached.
- FLP result still applies:
 - Paxos is **not guaranteed** to reach a consensus (ever or within a bounded time)

Paxos Problems

- Basic Paxos solves the problem for a single value
 - However, non-trivial to extend to Multi-Paxos. No agreement on the details of Multi-Paxos
 - We will discuss Multi-Paxos after we complete our discussion of Raft
- Doesn't fully address liveness
- Does not discuss cluster membership management

Led John Ousterhout and his PhD student Diego Ongara, to design a new consensus algorithm with **understandability as a primary design goal**

Next Lecture

- Raft: In Search of an Understandable Consensus Algorithm