

CS 582: Distributed Systems

Multi-Paxos



Dr. Zafar Ayyub Qazi

Fall 2024

Recap: Basic Paxos

- One or more replicas propose values
- All replicas must agree on a single value as chosen
- Only one value is ever chosen

Recap: Two-Phase Approach in Basic Paxos

- Phase 1: Broadcast **Prepare** RPC
 - Find out about any chosen values
 - Block older proposals that have not yet been completed
- Phase 2: Broadcast **Accept** RPC
 - Ask acceptors to accept a specific value

Recap: Basic Paxos

Proposers

- 1) Choose new proposal number n
- 2) Broadcast `Prepare(n)` to all servers
- 4) When responses received from majority:
 - If any `acceptedValues` returned, replace value with `acceptedValue` for highest `acceptedProposal`
- 5) Broadcast `Accept(n , value)` to all servers
- 6) When responses received from majority:
 - Any rejections (`result > n`)? goto (1)
 - Otherwise, **value is chosen**

Acceptors




- 3) Respond to `Prepare(n)`:
 - If $n > \text{minProposal}$ then $\text{minProposal} = n$
 - `Return(acceptedProposal, acceptedValue)`
- 6) Respond to `Accept(n , value)`:
 - If $n \geq \text{minProposal}$ then
 $\text{acceptedProposal} = \text{minProposal} = n$
 $\text{acceptedValue} = \text{value}$
 - `Return(minProposal)`

Acceptors must record `minProposal`, `acceptedProposal`, and `acceptedValue` on stable storage (disk)

Recap: Goal \rightarrow Replicated Log



Clients

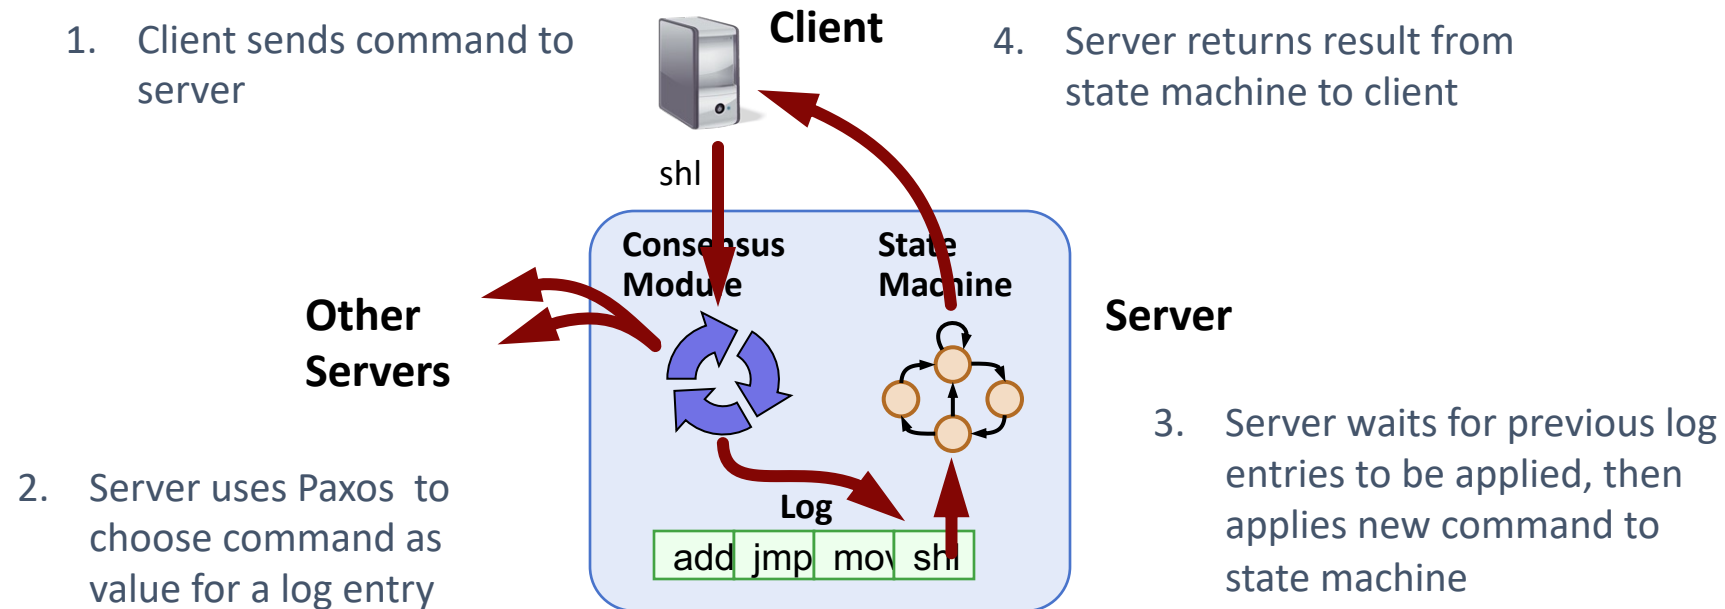


Servers

**Basic Paxos solves the problem for only a single log entry.
How do we extend it to solve the log replication problem?**

Multi-Paxos

- Separate instance of Basic Paxos for each entry in the log:
 - Add **index** argument to Prepare and Accept (selects entry in log)



Multi-Paxos Issues

- Which log entry to use for a given client request?
- Performance optimizations:
 - Use leader to reduce proposer conflicts
 - Eliminate most Prepare requests
- Ensuring full replication
- Client protocol
- Configuration changes

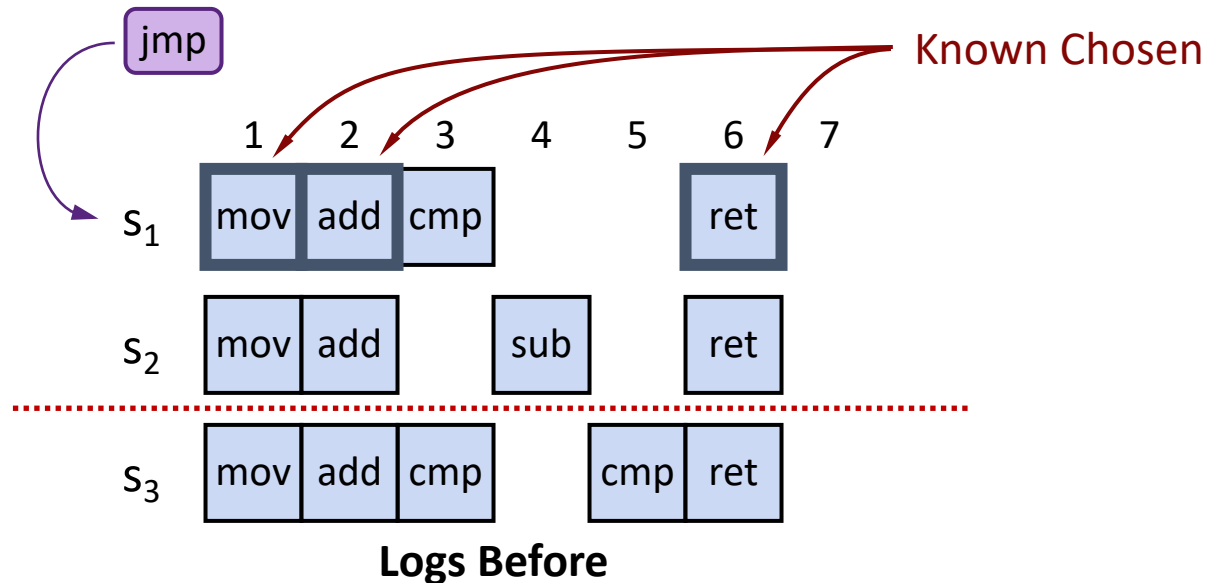
Note: Multi-Paxos not specified precisely in literature

Selecting Log Entries

- When request arrives from client:

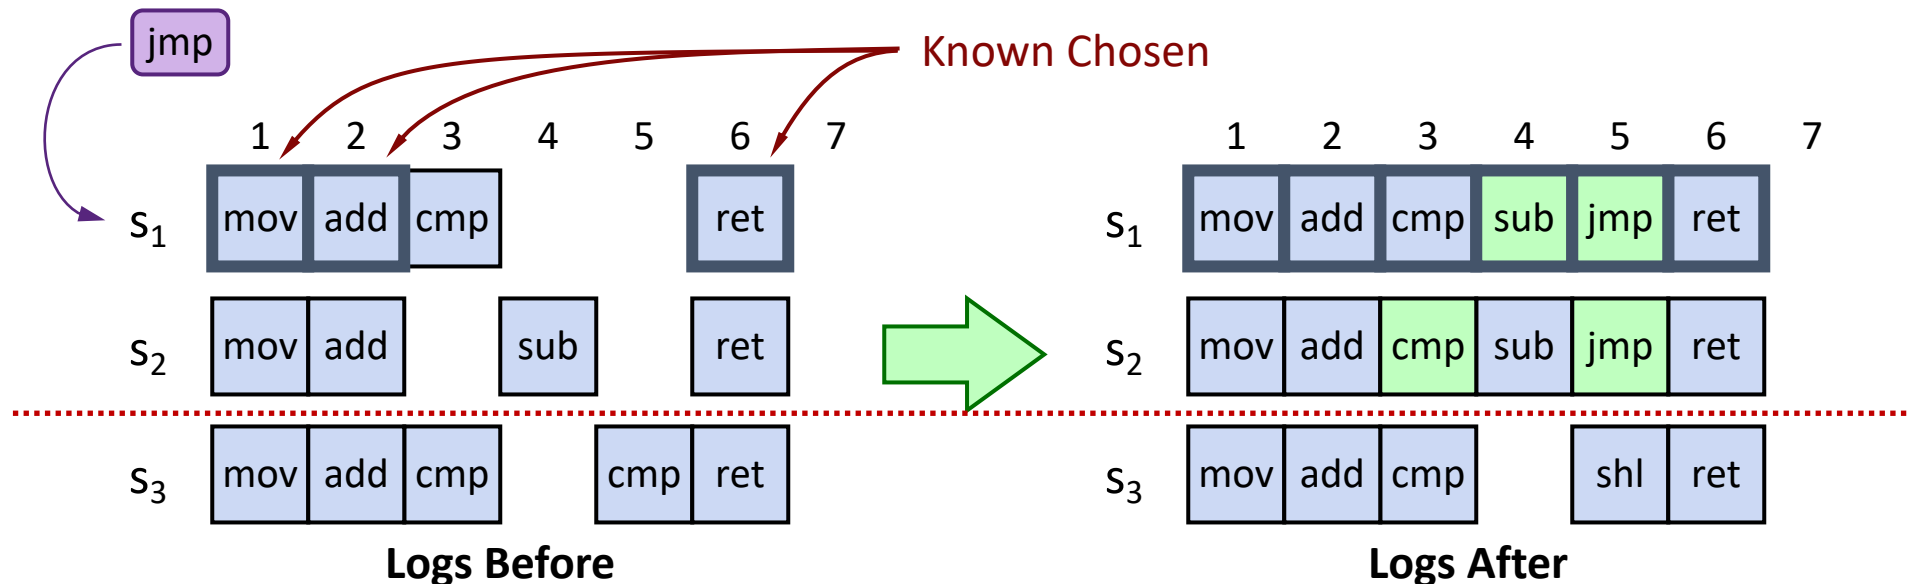
Selecting Log Entries

- When request arrives from client:
 - Find first log entry not known to be chosen



Selecting Log Entries

- When request arrives from client:
 - Find first log entry not known to be chosen
 - Run Basic Paxos to propose client's command for this index
 - Prepare returns acceptedValue?
 - Yes: finish choosing acceptedValue, start again
 - No: choose client's command



Selecting Log Entries (Cont'd)

- Servers can handle multiple client requests concurrently:
 - Select different log entries for each
- Must apply commands to state machine in log order

Improving Efficiency

Improving efficiency is a key goal for many organizations, as it allows them to do more with less. This can be achieved through a variety of methods, including process optimization, technology adoption, and resource management.

One of the most effective ways to improve efficiency is by streamlining processes. This involves identifying areas where time and resources are wasted and finding ways to eliminate or reduce those wastes. This can be done through a variety of methods, including process mapping, value stream mapping, and lean manufacturing.

Another way to improve efficiency is by adopting new technologies. This can include anything from automation to artificial intelligence. By using technology to automate repetitive tasks, organizations can free up their employees to focus on more complex, high-value work.

Finally, improving efficiency can also be achieved through better resource management. This involves ensuring that resources are allocated effectively and that there are no unnecessary costs. This can be done through a variety of methods, including budgeting, cost control, and performance management.

By implementing these strategies, organizations can improve their efficiency and achieve their goals more effectively. This can lead to increased productivity, reduced costs, and improved customer satisfaction.

Improving efficiency is a continuous process, and organizations should regularly review their processes and resources to ensure they are always operating at the highest level of efficiency.

By focusing on efficiency, organizations can ensure they are always moving forward and achieving their goals.

Improving efficiency is a key goal for many organizations, as it allows them to do more with less. This can be achieved through a variety of methods, including process optimization, technology adoption, and resource management.

Improving Efficiency

- Using Basic Paxos is inefficient:
 - With multiple concurrent proposers, **conflicts** and restarts are likely (higher load → more conflicts)
 - **2 rounds** of RPCs for each value chosen (Prepare, Accept)

Solution:

1. Pick a leader

- At any given time, only one server acts as Proposer

2. Eliminate most Prepare RPCs

- Prepare once for the entire log (not once per entry)
- Most log entries can be chosen in a single round of RPCs

Leader Election

One simple approach from Lamport:

- Let the **server with highest ID** act as leader
- **Each server sends a heartbeat message to every other server every T ms**
- If a server hasn't received heartbeat from server with higher ID in last $2T$ ms, it acts as leader:
 - Accepts requests from clients
 - Acts as proposer and acceptor
- **If server not leader:**
 - Rejects client requests (redirect to leader)
 - Acts only as acceptor

Eliminating Prepares

- Why is Prepare needed?
 - Block old proposals
 - Make proposal numbers refer to the **entire log**, not just one entry
 - Find out about (possibly) chosen values
 - Return highest proposal accepted for current entry
 - Also return **noMoreAccepted**: no proposals accepted for any log entry beyond current one
- If acceptor responds to Prepare with noMoreAccepted, skip future Prepares with that acceptor (until Accept rejected)
- Once leader receives noMoreAccepted from majority of acceptors, no need for Prepare RPCs
 - Only 1 round of RPCs needed per log entry (Accepts)

Full Disclosure

- So far, information flow is incomplete:
 - Log entries not fully replicated (majority only)
Goal: full replication
 - Only proposer knows when entry is chosen
Goal: all servers know about chosen entries
- Solution part 1/4: keep retrying Accept RPCs until all acceptors respond (in background)
 - Fully replicates most entries
- Solution part 2/4: track chosen entries
 - **Mark entries** that are known to be chosen:
 $\text{acceptedProposal}[i] = \infty$
 - Each server maintains **firstUnchosenIndex**: index of earliest log entry not marked as chosen

Full Disclosure (Cont'd)

- Solution part 3/4: proposer tells acceptors about chosen entries
 - Proposer includes its firstUnchosenIndex in Accept RPCs.
 - Acceptor marks all entries i chosen if:
 - $i < \text{request.firstUnchosenIndex}$
 - $\text{acceptedProposal}[i] == \text{request.proposal}$
 - Result: acceptors know about *most* chosen entries

| | | | | | | | | | |
|---|----------|----------|----------|-----|----------|----------|---|-----|---|
| log index | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
| acceptedProposal | ∞ | ∞ | ∞ | 2.5 | ∞ | 3.4 | | | |
| ... Accept(proposal = 3.4, index=8, value = v, firstUnchosenIndex = 7) ... | | | | | | | | | |
| | ∞ | ∞ | ∞ | 2.5 | ∞ | ∞ | | 3.4 | |

before Accept

after Accept

Still don't have complete information

Full Disclosure (Cont'd)

- Solution part 4/4: entries from old leaders
 - Acceptor returns its firstUnchosenIndex in Accept replies
 - If proposer's firstUnchosenIndex > firstUnchosenIndex from response, then proposer sends **Success** RPC (in background)
- Success(index, v): notifies acceptor of chosen entry:
 - acceptedValue[index] = v
 - acceptedProposal[index] = ∞
 - return its new firstUnchosenIndex
 - Proposer sends additional Success RPCs, if needed

Client Protocol (Similar to Raft)

- Send commands to leader
 - If leader unknown, contact any server
 - If contacted server not leader, it will redirect to leader
- Leader does not respond until command has been chosen for log entry and executed by leader's state machine
- If request times out (e.g., leader crash):
 - Client reissues command to some other server
 - Eventually redirected to new leader
 - Retry request with new leader

Client Protocol (Cont'd)


- What if the leader crashes after executing the command but before responding?
 - Must not execute command twice
- Client embeds a unique id in each command (just like Raft)
 - Server includes id in log entry
 - State machine records most recent command executed for each client
 - Before executing command, state machine checks to see if the command already executed, if so:
 - Ignore command
 - Return response from old command

Configuration Changes

- Consensus mechanism must support changes in the configuration:
 - Replace failed machine
 - Change degree of replication


Configuration Changes (Cont'd)

- Safety requirement:
 - During configuration changes, it must not be possible for different majorities to choose different values for the same log entry:



Configuration Changes (Cont'd)

- Paxos solution: use the log to manage configuration changes:
 - Configuration is stored as a log entry
 - Replicated just like any other log entry
 - Configuration for choosing entry i determined by entry $i-\alpha$.
Suppose $\alpha = 3$:



- Notes:
 - α limits concurrency: can't choose entry $i+\alpha$ until entry i chosen
 - Issue no-op commands if needed to complete change quickly

Paxos Summary

- Basic Paxos:
 - Prepare phase
 - Accept phase
- Multi-Paxos:
 - Choosing log entries
 - Leader election
 - Eliminating most Prepare requests
 - Full information propagation
- Client protocol
- Configuration changes

So far, we have assumed **fail-stop** failures
when dealing with consensus

Consensus with Fail-stop failures

- **Failure-Stop failures:** Nodes fail by crashing
 - A machine is either working correctly or it is doing nothing
- Paxos/Raft can ensure safety and solve the state machine replication problem as long as at least a **majority of nodes are up**
 - With $N = 2f + 1$ replicas can tolerate upto **f simultaneous fail-stop failures**
 - Ensure replicas execute operations in the **same order**
- **Remember FLP proof still applies:** can't sure both safety and liveness in an asynchronous system
 - Raft/Paxos provide liveness when at least a majority of nodes can communicate with reasonable timeliness

Byzantine Faults

- Nodes fail arbitrarily, and may
 - Perform incorrect computation
 - Send different and wrong messages
 - Not send message at all
 - Lie about the input value
 - Collude with other failed nodes
 - And more ...

Why care about Byzantine faults?

- Can be caused by
 - Malicious attack → increasing
 - Software errors → commonplace
- Example applications
 - Aircraft flight control systems (e.g., Boeing)
 - Blockchain platforms (e.g., Zilliqa, Hyperledger)
 - Byzantine-tolerant Distributed File System

Mini-case study: Boeing 737 Max

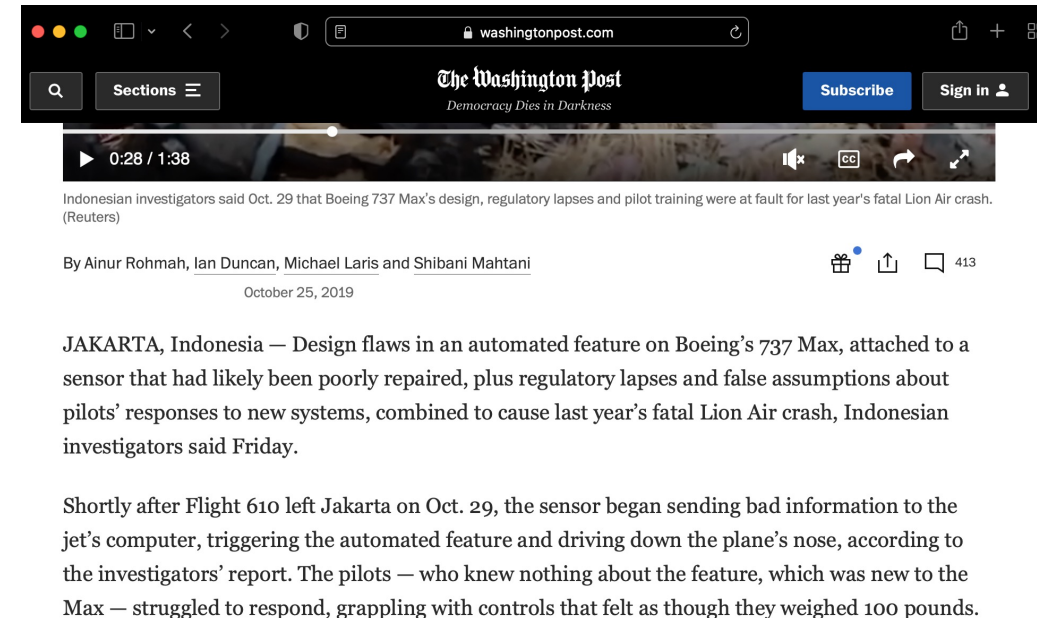
- Launched at the end of 2017
 - Had more efficient engines than 737 aircrafts
 - Introduced a new software system for automatic maneuvering of the plane (called MCAS)
- Multiple crashes
 - Lion Air Flight 610 – Oct 2018
 - Killing all 189 people on board
 - Ethiopian Airlines Flight 302 – March 2019
 - Killing all 157 people on board



Mini-case study: Boeing 737 Max (cont'd)

“Shortly after Flight 610 left Jakarta on Oct. 29, the sensor began sending bad information to the jet’s computer, triggering the automated feature and driving down the plane’s nose, according to the investigators’ report ...”

Washington Post




The screenshot shows a web browser displaying a Washington Post article. The browser's address bar shows 'washingtonpost.com'. The page header includes the Washington Post logo, the tagline 'Democracy Dies in Darkness', and buttons for 'Subscribe' and 'Sign in'. Below the header is a video player with a progress bar at 0:28 / 1:38. The article text begins with 'Indonesian investigators said Oct. 29 that Boeing 737 Max's design, regulatory lapses and pilot training were at fault for last year's fatal Lion Air crash. (Reuters)'. The byline reads 'By Ainur Rohmah, Ian Duncan, Michael Laris and Shibani Mahtani' and the date is 'October 25, 2019'. The main text of the article starts with 'JAKARTA, Indonesia — Design flaws in an automated feature on Boeing's 737 Max, attached to a sensor that had likely been poorly repaired, plus regulatory lapses and false assumptions about pilots' responses to new systems, combined to cause last year's fatal Lion Air crash, Indonesian investigators said Friday.' A second paragraph follows, repeating the information from the quote on the left: 'Shortly after Flight 610 left Jakarta on Oct. 29, the sensor began sending bad information to the jet's computer, triggering the automated feature and driving down the plane's nose, according to the investigators' report. The pilots — who knew nothing about the feature, which was new to the Max — struggled to respond, grappling with controls that felt as though they weighed 100 pounds.'


Byzantine Fault Tolerance



Design services that continue to function correctly despite Byzantine faults



Solve the replicated state machine problem machine with Byzantine faults



Solving consensus with Byzantine faults

Can we reach consensus in the presence of
Byzantine Faults with Paxos/Raft?

Can't use Paxos/Raft with Byzantine Faults


- The intersection of two majority ($f + 1$ node) quorums may be a **byzantine node**
- Byzantine node tells different quorums different things!
 - Leading to nodes committing different values
- Raft: **Can't rely on the leader to assign log index**
 - Could assign same log index to different requests

Bare majority quorums may not be enough in the presence of byzantine faults

A leader might be a byzantine node so we can't trust it

Byzantine Fault Tolerance

- Practical Byzantine Fault Tolerance (PBFT) Algorithm
 - [Liskov and Castro, 1999]



Barbara Liskov