

CS 582: Distributed Systems

PBFT and ZooKeeper



Dr. Zafar Ayyub Qazi

Fall 2024

Agenda

- Wrap-up discussion of PBFT
- ZooKeeper

Specific learning outcomes

By the end of today's lecture, you should be able to:

- ☐ Describe the core components and workflow of the PBFT algorithm
- ☐ Analyze the effectiveness of PBFT in maintaining consensus under Byzantine conditions
- ☐ Evaluate the performance of PBFT
- ☐ Explain the core design of ZooKeeper
- ☐ Analyze how ZooKeeper is able to scale system throughput with increasing replicas
- ☐ Analyze the consistency guarantees that ZooKeeper is providing


Byzantine Fault Tolerance



Design services that continue to function correctly despite Byzantine faults



Solve the replicated state machine problem with Byzantine faults



Solving consensus with Byzantine faults


Recap: Can't use Paxos/Raft with Byzantine Faults

Bare majority quorums may not be enough in the presence of byzantine faults

A leader might be a byzantine node so we can't trust it

Byzantine Fault Tolerance

- Practical Byzantine Fault Tolerance (PBFT) Algorithm
 - [Liskov and Castro, 1999]




Barbara Liskov

PBFT: Impact

- Seminal work on byzantine fault tolerance
 - Byzantine fault tolerance was for long considered an exotic topic
- PBFT showed Byzantine fault tolerance is possible under certain assumptions
 - Has inspired a large body of work in the last two decades
- Used by multiple blockchains like Zilliqa, Hyperledger fabric, etc.

PBFT Overview

- $3f + 1$ replicas to survive f byzantine failures
 - Shown to be minimal
- Quorum of $2f + 1$ replicas
- Three phases (instead of two)
- Primary-backup protocol
 - Since primary can be byzantine node, replicas observe, can trigger change
 - Change through the idea of views; primary = view mod IRI
- Clients: need $f + 1$ matching responses from different replicas



Key assumptions

- No more than f out of $3f + 1$ replicas can be faulty
 - The other $2f + 1$ replicas operate correctly – follow the PBFT protocol
- No client failure – clients can never do anything bad
- Worst case
 - A single attacker controlling f faulty replicas to break the system
- Note: faulty \rightarrow could be experiencing byzantine faults

What are the attacker's powers?

What are the attacker's powers?

- Supplies the code that faulty replicas run
- Knows the code the non-faulty replicas are running
- Knows the faulty replicas' crypto keys
- Can read network messages
- However, can't forge messages of non-faulty nodes
 - Messages are encrypted -- no guessing of crypto keys or breaking of cryptography

PBFT is a primary-backup protocol

- What can go wrong if we have a Byzantine primary?
 - It can send a wrong result to the client
 - Different updates to different replicas
 - Ignore a client request
- How do clients interact with the system?
 - If they just contact the primary, and the primary is byzantine, then the system is in trouble

Views


Views

- Identify each replica by an integer
 - For a set of R replicas $\{0, \dots, |R|-1\}$
- The replicas move through succession of configurations called **views**
- In a view, one replica is the primary and the rest are backups
- The primary of a view is a replica p such that $p = v \bmod |R|$ where v is the view number

How Clients determine success?

How Clients determine success?

- Clients wait for $f + 1$ identical replies from different replicas
- But \geq one reply is from a **non-faulty replica**



What Clients exactly do?

- Send requests to the **primary** replica
 - The primary multicasts the request to the backups
 - Replicas execute the request and send a reply to the client
- If the client does not receive replies soon enough, **it broadcasts the request to all replicas**
 - If the request has already been processed, the replicas simply resend reply
 - Replicas remember the last reply message they sent to each client
 - Otherwise, if the replica is not the primary it relays the request to the primary
 - If the primary does not multicast the request to the group, it will eventually be suspected to be faulty by enough replicas to cause a view change


What Replicas do?

- Carry out a protocol that ensures that
 - Replies from honest (non-faulty) replicas are correct
 - Enough replicas process each request to ensure that
 - The **non-faulty replicas** process the **same requests** and
 - In the **same order**
- Non-faulty replicas obey the protocol

Ordering Requests

Ordering Requests

- Primary picks the ordering of requests
 - But the primary might be a **byzantine node**




- Backups ensure primary behaves correctly
 - Check and certify ordering
 - Trigger view changes to replace faulty primary

Byzantine Quorums

Byzantine Quorums

- A Byzantine quorum contains $\geq 2f+1$ replicas (given $3f+1$ total nodes)



- One operation's quorum **overlaps** with the next operation's quorum
 - There are $3f+1$ replicas, in total
 - So overlap is $\geq f+1$ replicas
- $f+1$ replicas must contain ≥ 1 non-faulty replica

Message Authentication?

Message Authentication

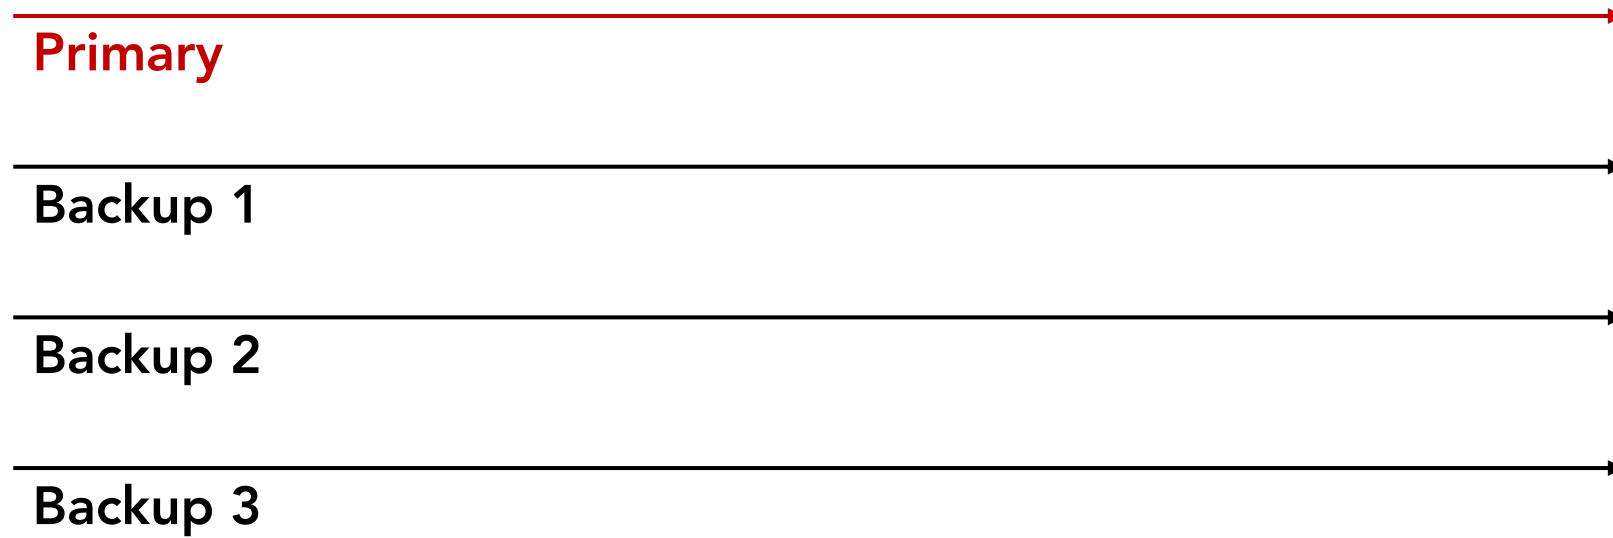
- Public-key cryptography for signatures
- Each client and server has a private and public key
- All hosts know all public keys
- Signed messages are signed with private key
- Public key can verify that message came from host with the private key

How many phases of interactions?


Three Phases

- **Pre-prepare:** pick order of request and inform replicas
- **Prepare:** ensures order within views
- **Commit:** ensures order across views

Normal Operation (primary is not faulty)



Normal Operation (primary is not faulty)



Normal Operation (primary is not faulty)

request:
 $m_{\text{Signed, Client}}$

Message
includes
operation,
timestamp,
client id

Assumption
A client sends
operations
one by one

Primary

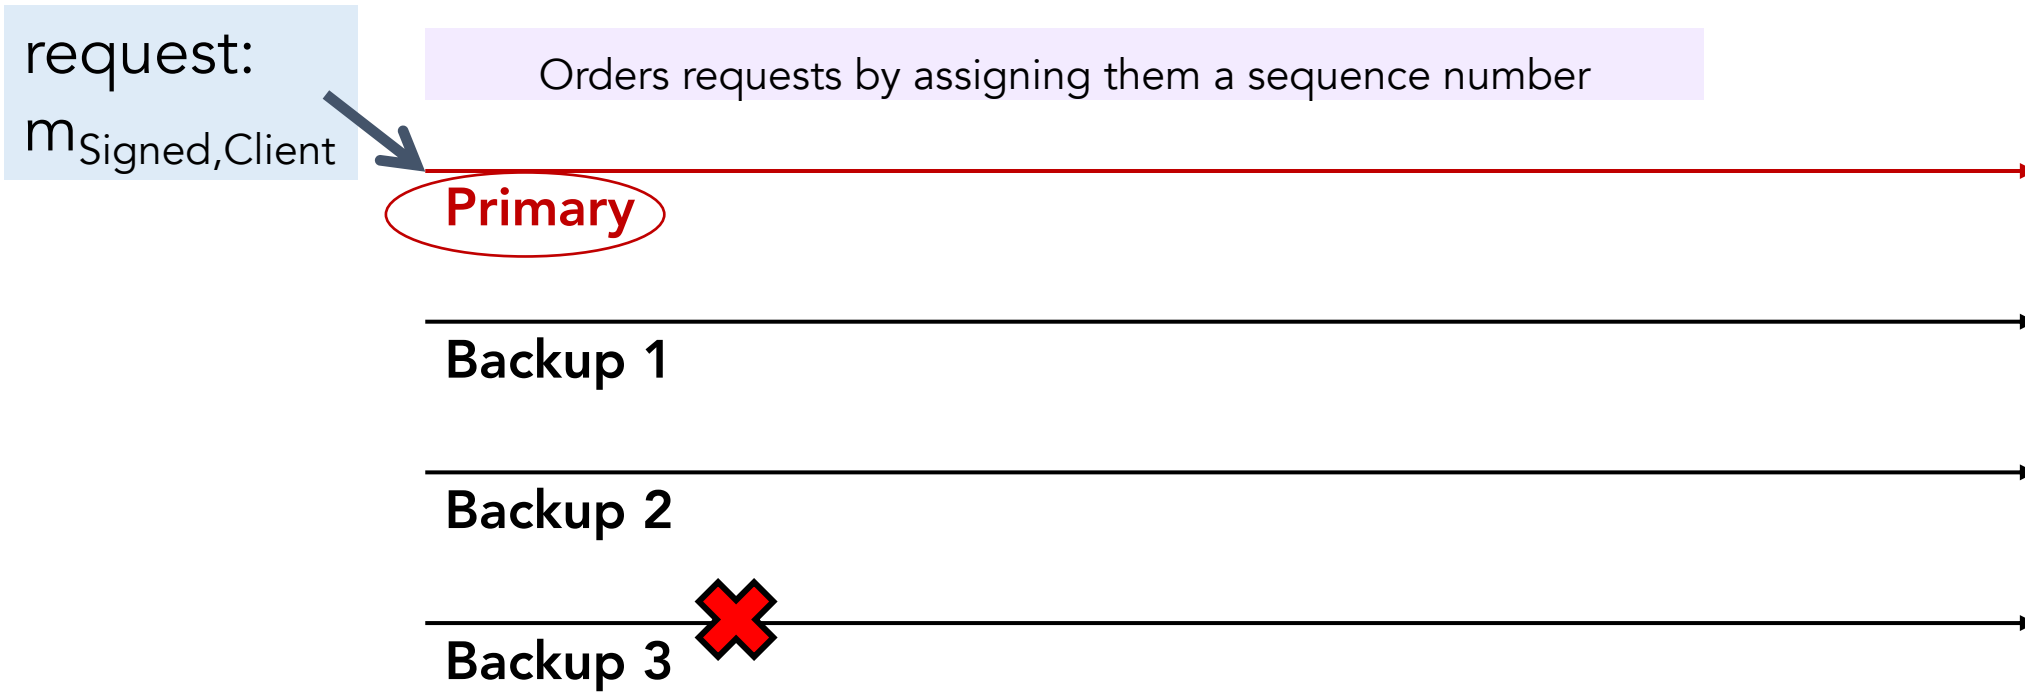
Backup 1

Backup 2

Backup 3

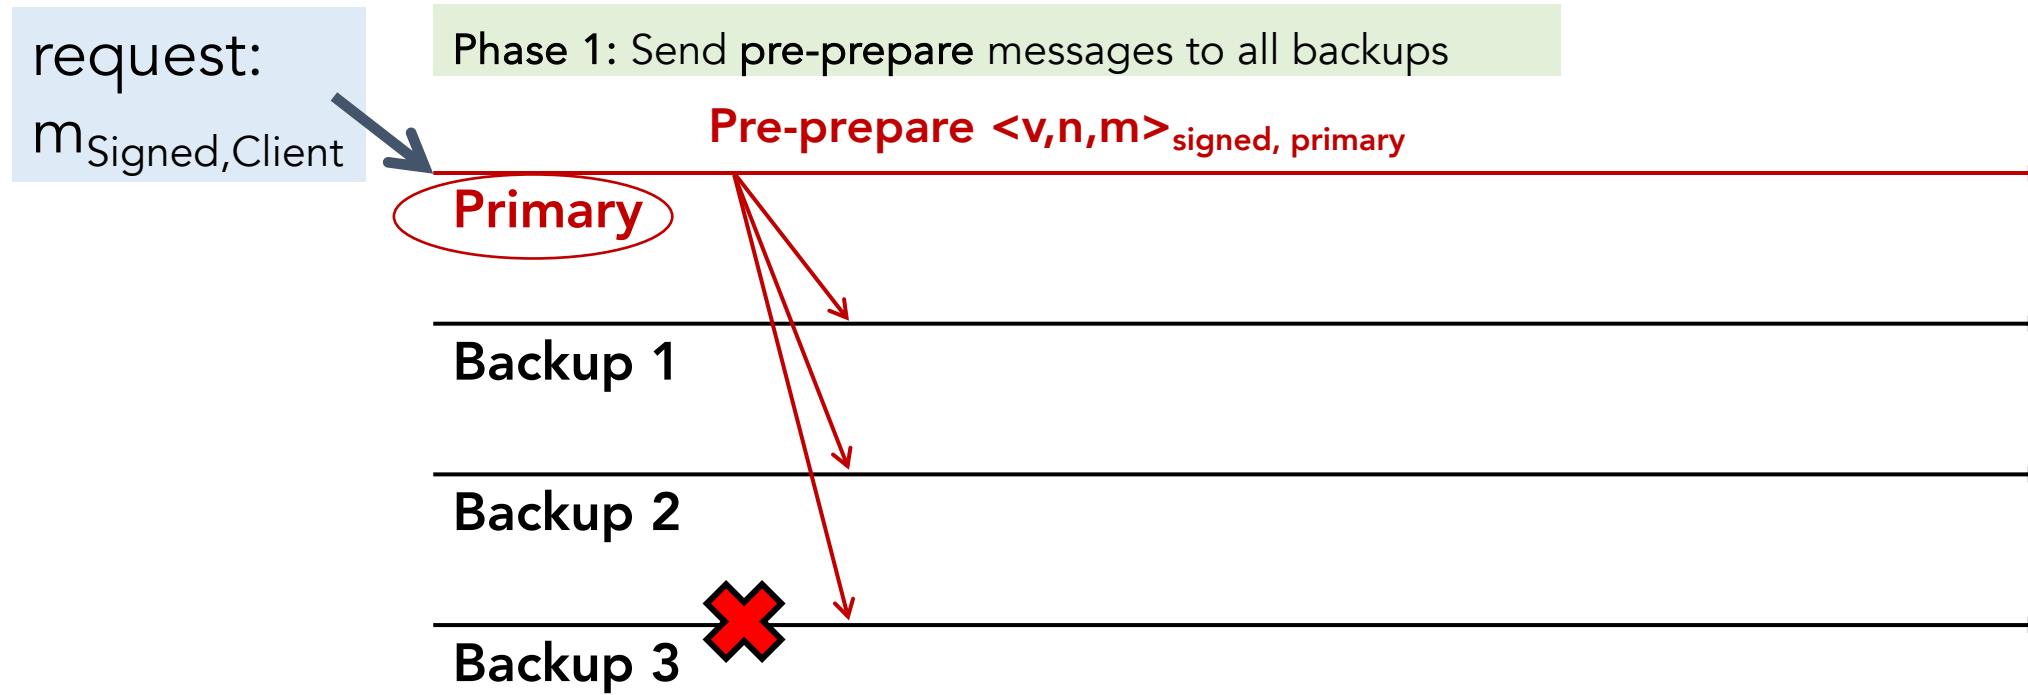


Normal Operation (primary is not faulty)

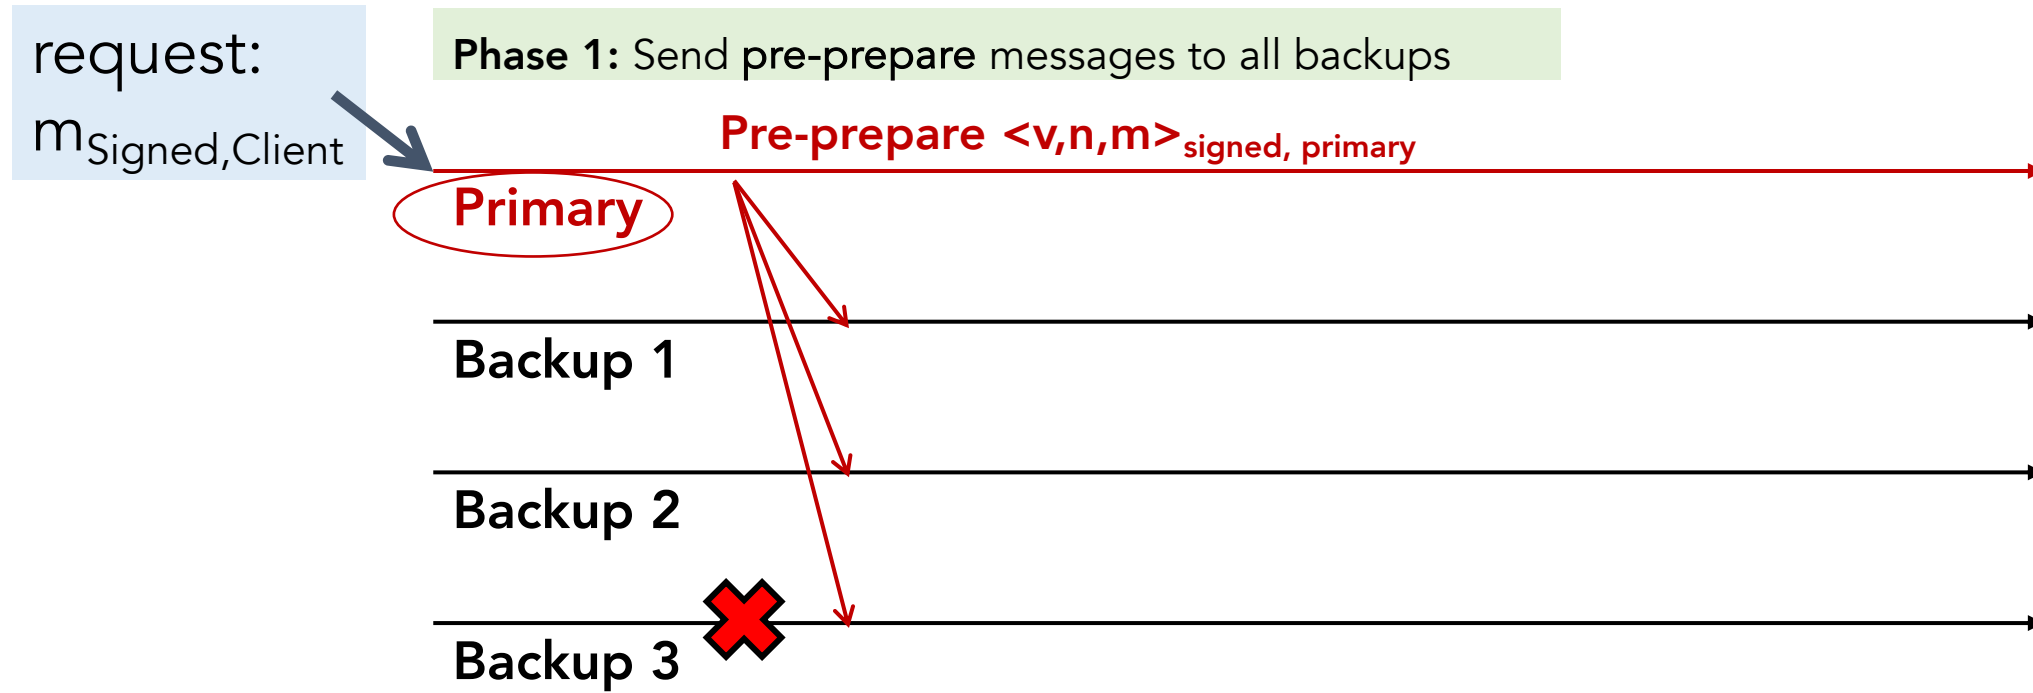


- Primary chooses the request's sequence number
 - Sequence number determines order of execution

Normal Operation (primary is not faulty)



Normal Operation (primary is not faulty)




- Backups do the following (and some other checks):
 - They check they are in view v
 - It has not seen another message with view v and sequence number n

Normal Operation (primary is not faulty)

request:

$m_{\text{Signed, Client}}$




Phase 2: Each backup which accepts pre-prepare msgs broadcasts a prepare message

Normal Operation (primary is not faulty)

request:

$m_{\text{Signed, Client}}$




Phase 2: Each backup which accepts pre-prepare msgs broadcasts a prepare message

- Backups wait for:
 - Prepared certificate: a collection of $2f$ matching prepare messages from replicas (including itself) + 1 matching pre-prepare message from the primary

Normal Operation (primary is not faulty)

request:

$m_{\text{Signed, Client}}$




Phase 2: Each backup which accepts pre-prepare msgs broadcasts a prepare message

- Backups wait for:
 - Prepared certificate: a collection of $2f$ matching prepare messages from replicas (including itself) + 1 matching pre-prepare message from the primary

Normal Operation (primary is not faulty)

request:

$m_{\text{Signed, Client}}$



Phase 2: Each backup which accepts pre-prepare msgs broadcasts a prepare message


- Backups wait for:
 - Pre Each **correct** node has a prepared certificate locally, from replicas (i.e. ≥ 1 correct nodes), but does not know whether the other correct nodes do too! So, we **can't commit** yet!

Normal Operation (primary is not faulty)

request:

$m_{\text{Signed,Client}}$


Phase 3: Each backup which has a **prepared** certificate broadcasts a **commit** message



Normal Operation (primary is not faulty)

request:

$m_{\text{Signed, Client}}$




Phase 3: Each backup which has a prepared certificate broadcasts a commit message

- Backups wait for:
 - **Commit certificate:** a collection of $2f+1$ matching commit messages
 - Same view, sequence number and message

Normal Operation (primary is not faulty)

request:

$m_{\text{Signed,Client}}$



Phase 3: Each backup which has a prepared certificate broadcasts a commit message

- Backups wait for:
 - Commit certificate: a collection of $2f+1$ matching commit messages
 - Since the primary is not faulty, it will send a commit message
- Once the request is **committed**, replicas execute the operation and send a reply directly back to the client.

Byzantine Primary



Primary

Backup 1

Backup 2

Backup 3

Byzantine Primary



Primary

Pre-prepare $\langle v, n, m \rangle$ signed, primary

Backup 1


Pre-prepare $\langle v, n, m' \rangle$ signed, primary

Backup 2

Pre-prepare $\langle v, n, m'' \rangle$ signed, primary

Backup 3

Byzantine Primary



Byzantine primary

- In general, backups **won't prepare** if primary lies
- **Suppose they did:** two distinct requests m and m' for the same sequence number n
 - Then prepared quorum certificates (each of size $2f + 1$) would **intersect** at an **honest** replica
 - So that honest replica would have sent a prepare message for both m and m'
 - **So $m = m'$**

View Change

- If a replica suspects the primary is faulty, it requests a *view change*
 - Sends a *view change* request to all replicas
 - Everyone acks the view change request
- New primary collects a quorum of $(2f+1)$ responses
 - Sends a *new-view* message with this certificate
- Need committed operations to survive into next view
 - Client may have gotten answer
 - View change request contain checkpoints + newer prepare certificates

Other Bits

- Garbage collection
 - Can't let log grow without bound
 - So shrink log when its gets too big
- Proactive recovery
 - Recover the replica to a known good state whether faulty or not

PBFT Performance

- How much time does it take to commit requests?
- At least 3 Round Trip Network Delays + Disk Writes

Failures can be correlated

- You have to be careful about the assumption of independent failures

Non-deterministic state machines

- Many services may have some form of non-determinism, e.g.,
 - Timers
 - Random numbers
- Paxos/Raft/PBFT by default designed for deterministic state machines

Summary of key ideas

- $2f+1$ Quorum
 - We need $2f+1$ replicas in a quorum to deal with byzantine faults
 - Assuming a total of $3f+1$ replicas with atmost f byzantine faults
- Primary backup with view changes
- Three Rounds
 - Pre-prepare: pick order of request and inform clients
 - Prepare: ensures order within views
 - Commit: ensures order across views
- Replicas directly contact the clients
 - Clients wait for $f+1$ matching responses

PBFT Conclusion

- Byzantine fault tolerance was for long considered an exotic topic
 - 1980s focused primarily on fail-stop failures
- PBFT showed Byzantine fault tolerance is possible under certain assumptions
- But there were still challenges around performance
 - Challenges: lots of coordination with three phases
 - The paper and a lot of the followup work is about making byzantine fault tolerance more performant

ZooKeeper

ZooKeeper

- General-purpose distributed coordination service
- Many use cases
 - Configuration management, leader election, naming service, data synchronization, etc.
- Used by many companies
 - Yahoo, Yelp, Reddit, Facebook, Twitter, eBay, Rackspace, etc.

Performance Question

- In a replicated system with N replica servers, can we get N times higher performance?

ZooKeeper Setup

ZooKeeper Performance

Servers	100% Reads	0% Reads
13	460k	8k
9	296k	12k
7	257k	14k
5	165k	18k
3	87k	21k

Table 1: The throughput performance of the extremes of a saturated system.