

CS 582: Distributed Systems

Scaling Distributed Machine Learning with Parameter Server



Dr. Zafar Ayyub Qazi
Fall 2024

Why this paper?

- Different use case: distributed machine learning
- Influential design
 - E.g., TensorFlow's distributed execution uses it
- Relaxed consistency
 - For many ML applications, some inconsistency is OK
- Impressive evaluation results
 - On peta bytes of real data with billions of parameters

Machine Learning Primer

- Models are function approximators
 - The true function is unknown, so we learn an approximation from the data
- Examples:
 - $f(\text{user profile}) \rightarrow$ likelihood of ad click
 - $f(\text{picture}) \rightarrow$ likelihood picture contains a cat
 - $f(\text{words in document}) \rightarrow$ topics/search terms for document

Machine Learning Primer (Cont'd)

- Two Phases: Training and Inference
- During training, expose the model to many examples of data
 - Supervised: uses labeled data
 - Unsupervised: uses unlabelled data
- During inference, apply the trained model to get predictions for unseen data

Training data can be in the order of PBs

**Parameter server is about making
the training phase efficient**

Model Parameters

- Machine learning models learn parameter values
 - Large models can have billions/trillions of parameters
 - GPT-4 has more than a trillion parameters
- Training iterates thousands of times to incrementally tune the parameters' values
 - Popular algorithm: gradient descent

Challenges

- Need many workers
 - Because training data are too large for one machine
 - For parallel speedup
 - Parameters may not fit on a single machine either
- All workers need access to parameters
 - Coordination overheads: network bandwidth cost & synchronization delays
- Fault Tolerance is critical at scale

High Level Approach

- Distribute parameters and training data over multiple machines
- Devise an efficient coordination mechanism
 - That doesn't consume too much network or cause large delays in coordination

Discussion


- How to partition parameters & data?
- Do worker nodes need all parameters?
- How do workers access shared parameters?

Discussion

- How to partition parameters & data?
 - Consistent hashing is used to partition parameters
 - Process in parallel different partitions
 - Training data is partitioned across workers by a scheduler node
- Do worker nodes need all parameters?
- How do workers access shared parameters?

Discussion

- How to partition parameters & data?
 - Consistent hashing is used to partition parameters
 - Process in parallel different partitions
 - Training data is partitioned across workers by a scheduler node
- Do worker nodes need all parameters?
 - A small fraction of parameters needed by each worker if the number of workers is large
- How do workers access shared parameters?



Discussion

- How to partition data?
 - Consistent hashing is used to partition parameters
 - Process in parallel different partitions
 - Training data is partitioned across workers by a scheduler node
- Do worker nodes need all parameters?
 - A small fraction of parameters needed by each worker if the number of workers is large
- How do workers access shared parameters?
 - When parameters get updated, other nodes may need to be informed
 - How to we update the parameters?

Strawman1: Broadcast Updates

- Broadcast parameter changes between workers at end of training iteration
 - Workers exchange their parameter changes and then apply them once all have received all changes
- Possible Problems?
 - All-to-All broadcasts can generate huge amounts of traffic
 - Need to wait for all other workers before proceeding → idle time

Strawman2: Use a Coordinator/Leader

- A single coordinator collects and distributes updates at end of the training iteration
 - Workers send their changes to the coordinator
 - The coordinator collects, aggregates, and sends aggregated updates to workers
 - Workers modify their local parameters
- Possible Problems?
 - Single coordinator gets congested with updates
 - Single coordinator is a single point of failure


Another Solution

- Use multiple coordinators and replicate the coordinators for fault tolerance
- How to replicate it?

Parameter Server

- Parameter Servers (PS) and Workers
 - Other components include resource manager, server manager, and task scheduler
- Parameter Servers aggregate updates from workers

High-Level Architecture



Process

- Initialize parameters at PS
- Push to workers on each iteration: assign training tasks to workers
- Workers compute parameter updates
- Workers push parameter updates to the responsible parameter servers

Process (Cont'd)

- Parameter servers update parameters via user-defined function
 - Possibly aggregating parameter changes from multiple workers
- Parameter servers replicate changes
 - Then ACK to worker
- Once done worker pulls a new parameter value

Key-Value Interface

- Parameters often abstracted as big vector $w[0, \dots, z]$ for z parameters
 - Each logical vector position stores (key, value) which can be indexed by key
- Applies operations (push/pull/updated) on key ranges, not single parameters
- Why?
 - Improved efficiency due to batching

Further Optimizations

- Skip unchanged parameters
- Skip keys with value zero in data for range
 - Can also use threshold to drop unimportant updates

Fault Tolerance

- What if a worker crashes?
- What if a parameter server crashes?

Fault Tolerance

- What if a worker crashes?
 - Restart on another machine; load training data, pull parameters, continue or just drop it -- will only lose a small part of training data set usually doesn't affect outcome much, or training may take a little longer
- What if a parameter server crashes?
 - Lose all parameters stored there
 - Use consistent hashing to replicate data
 - On failures, neighboring backup takes over

Relaxed Consistency

- Many ML algorithms tolerate somewhat stale parameters in training
- **Intuition**: if parameters only change a little, not too bad to use old ones won't go drastically wrong (e.g., cat likelihood 85% instead of 91%)
 - Still converges to a decent model though may take longer (more training iterations due to high errors)
 - And more resource-efficient

Vector Clocks

- Need a mechanism to synchronize
 - When strong consistency is needed
 - Even with relaxed consistency when some workers may be very slow
 - Avoid some parameters getting very stale
- Workers need to be aware of how far along others and the servers are
- For this purpose, vector clocks are used

Vector Clocks (Cont'd)

- But vector clock for each key won't scale
- Vector clocks for ranges of keys (as Parameter Server uses ranges)

Summary

- Influential design – impacted the distributed deep learning frameworks
- Contribution: Synthesizes several existing techniques in a different (new) context
 - Data partitioning via Consistent hashing
 - Replication via Consistent hashing
 - Vector Clocks for synchronization
 - Flexible and Relaxed Consistency